

To-do list

1. Section X naive classification	4
2. Section X	8

Contents

5 I-priors for categorical responses	2
5.1 A latent variable motivation: the I-probit model	4
5.2 Identifiability and IIA	7
5.3 Estimation	10
5.3.1 Laplace approximation	11
5.3.2 Variational EM algorithm	13
5.3.3 Markov chain Monte Carlo methods	14
5.3.4 Comparison of estimation methods	15
5.4 A variational algorithm	17
5.4.1 Latent propensities \mathbf{y}^*	20
5.4.2 I-prior random effects \mathbf{w}	21
5.4.3 Kernel parameters η	22
5.4.4 Intercepts $\boldsymbol{\alpha}$	23
5.4.5 The CAVI algorithm	23
5.5 Post-estimation	24
5.6 Computational consideration	24
5.7 Examples	24
5.8 Conclusion	24
5.9 Miscellanea	27
5.10 Derivation of the CAVI algorithm	27
5.10.1 Derivation of $\tilde{q}(\mathbf{y}^*)$	29
5.10.2 Derivation of $\tilde{q}(\mathbf{w})$	30

5.10.3	Derivation of $\tilde{q}(\eta)$	33
5.10.4	Derivation of $\tilde{q}(\Psi)$	36
5.10.5	Derivation of $\tilde{q}(\alpha)$	38
Bibliography		40

Haziq Jamil

Department of Statistics

London School of Economics and Political Science

PhD thesis: ‘Regression modelling using Fisher information covariance kernels (I-priors)’

Chapter 5

I-priors for categorical responses

chapter5

In a regression setting such as (1.1), consider polytomous response variables y_1, \dots, y_n , where each y_i takes on exactly one of the values from the set of m possible choices $\mathcal{M} = \{1, \dots, m\}$. Modelling categorical response variables is of profound interest in statistics, econometrics and machine learning, with applications aplenty. In the social sciences, categorical variables often arise from survey responses, and one may be interested in studying correlations between explanatory variables and the categorical response of interest. Economists are often interested in discrete choice models to explain and predict choices between several alternatives, such as consumer choice of goods or modes of transport. In this age of big data, machine learning algorithms are used for classification of observations based on what is usually a large set of variables or features.

The normality assumption (1.2) is not entirely appropriate anymore. As an extension to the I-prior methodology, we propose a flexible modelling framework suitable for regression of categorical response variables. In the spirit of generalised linear models (McCullagh and Nelder, 1989), we relate class probabilities of the observations to a normal I-prior regression model via a link function. Perhaps though, it is more intuitive to view it as machine learners do: since the regression function is ranged on the entire real line, it is necessary to “squash” it through some sigmoid function to conform it to the interval $[0, 1]$ suitable for probability ranges.

Expanding on this idea further, assume that the y_i ’s follow a categorical distribution, denoted by

$$y_i \sim \text{Cat}(p_{i1}, \dots, p_{im}),$$

with the class probabilities satisfying $p_{ij} \geq 0, \forall j = 1, \dots, m$ and $\sum_{j=1}^m p_{ij} = 1$. The probability mass function (pmf) of y_i is given by

$$p(y_i) = p_{i1}^{[y_i=1]} \dots p_{im}^{[y_i=m]}$$

where the notation $[\cdot]$ refers to the Iverson bracket¹. The dependence of the class probabilities on the covariates is specified through the relationship

$$g(p_{i1}, \dots, p_{im}) = (\alpha_1 + f_1(x_i), \dots, \alpha_m + f_m(x_i))$$

where $g : [0, 1]^m \rightarrow \mathbb{R}^m$ is some specified link function. As we will see later, a normal regression model as in (1.1) subject to (1.2) naturally implies a *probit* link function. With an I-prior assumed on the f_j 's, we call this method of probit regression using I-priors the *I-probit* regression model.

Due to the nature of the model, unfortunately, the posterior distribution of the regression functions cannot be found in closed form. In particular, marginalising the I-prior from the joint likelihood involves a high-dimensional intractable integral. We explore a fully Bayesian approach to estimate I-probit models using *variational inference*. The main idea is to replace the difficult posterior distribution with an approximation that is tractable. Working in a Bayesian setting together with variational inference allows us to estimate the model much faster than traditional MCMC sampling methods, yet provides us with the conveniences that come with Bayesian machinery. For example, inferences around log-odds is usually cumbersome for probit models, but a credibility interval can easily be obtained by resampling methods from the relevant posteriors, which are typically made up of densities which are familiar and readily available in software.

By choosing appropriate RKHSs/RKKSs for the regression functions, we are able to fit a multitude of binary and multinomial models, including multilevel or random-effects models, linear and non-linear classification models, and even spatio-temporal models. Examples of these models applied to real-world data is shown in Section 5.7. We find that the many advantages of the normal I-prior methodology transfer over quite well to the I-probit model for binary and multinomial regression.

¹ $[A]$ returns 1 if the proposition A is true, and 0 otherwise. The Iverson bracket is a generalisation of the Kronecker delta.

5.1 A latent variable motivation: the I-probit model

It is convenient, as we did in [Section X naive classification](#), to again think of the responses $y_i \in \{1, \dots, m\} = \mathcal{M}$ as comprising of a binary vector $\mathbf{y}_i = (y_{i1}, \dots, y_{im})^\top$, with a single ‘1’ at the position corresponding to the value that y_i takes. That is,

$$y_{ik} = \begin{cases} 1 & \text{if } y_i = k \\ 0 & \text{if } y_i \neq k. \end{cases}$$

With $y_i \stackrel{\text{iid}}{\sim} \text{Cat}(p_{i1}, \dots, p_{im})$ for $i = 1, \dots, n$, each y_{ij} is distributed as Bernoulli with probability p_{ij} , $j = 1, \dots, m$ according to the above formulation. Now, assume that, for each y_{i1}, \dots, y_{im} , there exists corresponding *continuous, underlying, latent variables* $y_{i1}^*, \dots, y_{im}^*$ such that

$$y_i = \begin{cases} 1 & \text{if } y_{i1}^* \geq y_{i2}^*, y_{i3}^*, \dots, y_{im}^* \\ 2 & \text{if } y_{i2}^* \geq y_{i1}^*, y_{i3}^*, \dots, y_{im}^* \\ \vdots & \\ m & \text{if } y_{im}^* \geq y_{i2}^*, y_{i3}^*, \dots, y_{i,m-1}^*. \end{cases} \quad (5.1)$$

{eq:latentmodel}

In other words, $y_{ij} = \arg \max_{k=1}^m y_{ik}^*$. Such a formulation is common in economic choice models, and is rationalised by a utility-maximisation argument: an agent faced with a choice from a set of alternatives will choose the one which benefits them most. In this sense, the y_{ij}^* ’s represent individual i ’s *latent propensities* for choosing alternative j .

Instead of modelling the observed y_{ij} ’s directly, we model instead the n latent variables in each class $j = 1, \dots, m$ according to the regression problem

$$y_{ij}^* = \alpha + \alpha_j + f_j(x_i) + \epsilon_{ij} \quad (5.2)$$

$$(\epsilon_{i1}, \dots, \epsilon_{im})^\top \stackrel{\text{iid}}{\sim} N_m(\mathbf{0}, \mathbf{\Psi}^{-1}).$$

{eq:multinomial-latent}

We can see some semblance of this model with the one in ??, and ultimately the aim is to assign I-priors to the regression function of these latent variables, which we shall describe shortly. For now, write $\boldsymbol{\mu}(x_i) \in \mathbb{R}^m$ whose j ’th component is $\alpha + \alpha_j + f_j(x_i)$, and realise that each $\mathbf{y}_i^* = (y_{i1}^*, \dots, y_{im}^*)^\top$ has the distribution $N_m(\boldsymbol{\mu}(x_i), \mathbf{\Psi}^{-1})$, conditional on the data x_i , the intercepts $\alpha, \alpha_1, \dots, \alpha_m$, the evaluations of the functions at x_i for each class $f_1(x_i), \dots, f_m(x_i)$, and the error covariance matrix $\mathbf{\Psi}^{-1}$.

The probability p_{ij} of observation i belonging to class j is calculated as

$$\begin{aligned} p_{ij} &= \mathbb{P}(y_i = j) \\ &= \mathbb{P}(\{y_{ij}^* > y_{ik}^* \mid \forall k \neq j\}) \\ &= \int \cdots \int_{\{y_{ij}^* > y_{ik}^* \mid \forall k \neq j\}} \phi(y_{i1}^*, \dots, y_{im}^* \mid \boldsymbol{\mu}(x_i), \boldsymbol{\Psi}^{-1}) dy_{i1}^* \cdots dy_{im}^*, \end{aligned} \quad (5.3)$$

{eq:p_{ij}}

where $\phi(\cdot \mid \boldsymbol{\mu}, \boldsymbol{\Sigma})$ is the density of the multivariate normal with mean $\boldsymbol{\mu}$ and variance $\boldsymbol{\Sigma}$. This is the probability that the normal random variable \mathbf{y}_i^* belongs to the set $\mathcal{C}_j := \{y_{ij}^* > y_{ik}^* \mid \forall k \neq j\}$, which are cones in \mathbb{R}^m . Since the union of these cones is the entire m -dimensional space of reals, the probabilities add up to one and hence they represent a proper probability mass function for the classes. While this does not have a closed-form expression and highlights one of the difficulties of working with probit models, the integral is by no means impossible to compute—see ?? for a note regarding this matter.

Now, we'll see how to specify an I-prior on the regression problem (5.2). In the naïve I-prior model, we wrote $f(x_i, j) = \alpha_j + f_j(x_i)$, and called for f to belong to an ANOVA RKKS with kernel defined in ?. Instead of doing the same, we take a different approach. Treat the α_j 's in (5.2) as intercept parameters to estimate with the additional requirement that $\sum_{j=1}^m \alpha_j = 0$. Further, let \mathcal{F} be a (centred) RKHS/RKKS of functions over \mathcal{X} with reproducing kernel h_η . Now, consider putting an I-prior on the regression functions $f_j \in \mathcal{F}$, $j = 1 \dots, m$, defined by

$$f_j(x_i) = f_0(x_i, j) + \sum_{k=1}^n h_\eta(x_i, x_k) w_{ik}$$

with $\mathbf{w}_i := (w_{i1}, \dots, w_{im})^\top \stackrel{\text{iid}}{\sim} \mathcal{N}(0, \boldsymbol{\Psi})$. This is similar to the naïve I-prior specification ??, except that the intercept have been treated as parameters rather than accounting for them using an RKHS of functions (Pearson RKHS or identity kernel RKHS). Importantly, the overall regression relationship still satisfies the ANOVA functional decomposition, because the α_j 's sum to zero. We find that this approach bodes well down the line computationally.

We call the multinomial probit regression model of (5.1) subject to (5.2) and I-priors on $f_j \in \mathcal{F}$, the *I-probit model*. For completeness, this is stated again: for $i = 1, \dots, n$,

$y_i = \arg \max_{k=1}^m y_{ik}^* \in \{1, \dots, m\}$, where, for $j = 1, \dots, m$,

$$\begin{aligned} y_{ij}^* &= \alpha + \alpha_j + \overbrace{f_0(x_i, j) + \sum_{k=1}^n h_\eta(x_i, x_k) w_{ik}}^{f_j(x_i)} + \epsilon_{ij} \\ \boldsymbol{\epsilon}_{i\cdot} &:= (\epsilon_{i1}, \dots, \epsilon_{im})^\top \stackrel{\text{iid}}{\sim} \text{N}_m(\mathbf{0}, \boldsymbol{\Psi}^{-1}) \\ \mathbf{w}_{i\cdot} &:= (w_{i1}, \dots, w_{im})^\top \stackrel{\text{iid}}{\sim} \text{N}_m(\mathbf{0}, \boldsymbol{\Psi}). \end{aligned} \tag{5.4}$$

{eq:iprobit
mod}

The parameters of the I-probit model are denoted by $\theta = \{\alpha_1, \dots, \alpha_m, \eta, \boldsymbol{\Psi}\}$. To establish notation, let

- $\boldsymbol{\epsilon} \in \mathbb{R}^{n \times m}$ denote the matrix containing (i, j) entries ϵ_{ij} , whose rows are $\boldsymbol{\epsilon}_{i\cdot}$, columns are $\boldsymbol{\epsilon}_{\cdot j}$, and is distributed $\boldsymbol{\epsilon} \sim \text{MN}_{n,m}(\mathbf{0}, \mathbf{I}_n, \boldsymbol{\Psi}^{-1})$;
- $\mathbf{w} \in \mathbb{R}^{n \times m}$ denote the matrix containing (i, j) entries w_{ij} , whose rows are $\mathbf{w}_{i\cdot}$, columns are $\mathbf{w}_{\cdot j}$, and is distributed $\mathbf{w} \sim \text{MN}_{n,m}(\mathbf{0}, \mathbf{I}_n, \boldsymbol{\Psi})$;
- $\mathbf{f}, \mathbf{f}_0 \in \mathbb{R}^{n \times m}$ denote the matrices containing (i, j) entries $f_j(x_i)$ and $f_0(x_i, j)$ respectively, so that $\mathbf{f} = \mathbf{f}_0 + \mathbf{H}_\eta \mathbf{w} \sim \text{MN}_{n,m}(\mathbf{1}_n \mathbf{f}_0^\top, \mathbf{H}_\eta^2, \boldsymbol{\Psi})$;
- $\boldsymbol{\alpha} = (\alpha + \alpha_1, \dots, \alpha + \alpha_m)^\top \in \mathbb{R}^m$ be the vector of intercepts;
- $\boldsymbol{\mu} = \mathbf{1}_n \boldsymbol{\alpha}^\top + \mathbf{f}$, whose (i, j) entries are $\mu_j(x_i) = \alpha + \alpha_j + f_j(x_i)$; and
- $\mathbf{y}^* \in \mathbb{R}^{n \times m}$ denote the matrix containing (i, j) entries y_{ij}^* , that is, $\mathbf{y}^* = \boldsymbol{\mu} + \boldsymbol{\epsilon}$, so $\mathbf{y}^* | \mathbf{w} \sim \text{MN}_{n,m}(\boldsymbol{\mu} = \mathbf{1}_n \boldsymbol{\alpha}^\top + \mathbf{H}_\eta \mathbf{w}, \mathbf{I}_n, \boldsymbol{\Psi}^{-1})$ and $\text{vec } \mathbf{y}^* \sim \text{N}_{nm}(\text{vec}(\mathbf{1}_n \boldsymbol{\alpha}^\top), \boldsymbol{\Psi} \otimes \mathbf{H}_\eta^2 + \boldsymbol{\Psi}^{-1} \otimes \mathbf{I}_n)$ (note that the marginal distribution of \mathbf{y}^* cannot be expressed as a matrix normal, except when $\boldsymbol{\Psi} = \mathbf{I}_m$).

Before proceeding with estimating the I-probit model (5.4), we lay out several standing assumptions:

A4 Centred responses. Set $\alpha = 0$.

A5 Zero prior mean. Assume a zero prior mean $f_0(x) = 0$ for all $x \in \mathcal{X}$.

A6 Fixed error precision. Assume $\boldsymbol{\Psi}$ is fixed.

Assumption A4 is a requirement for identifiability, while A5 is motivated by a similar argument to assumption A2 in the normal I-prior model. While estimation of $\boldsymbol{\Psi}$ would add flexibility to the model, several computational issues were not able to be resolved within the time limitations of completing this project (see ??).

ass:A4

ass:A5

ass:A6

sec:ia

5.2 Identifiability and IIA

The parameters in the standard linear multinomial probit model is well known to be unidentified (Keane, 1992; Train, 2009), and we find this to be the case in the I-probit model as well. Unrestricted probit models are not identified for two reasons. Firstly, an addition of a non-zero constant $a \in \mathbb{R}$ to the latent variables y_{ij}^* 's in (5.1) will not change which latent variable is maximal, and therefore leaves the model unchanged. It is for this reason assumptions A4 and A5 are imposed. Secondly, all latent variables can be scaled by some positive constant $c \in \mathbb{R}_{>0}$ without changing which latent variable is largest. This means that m -variate normal distribution $N_m(\boldsymbol{\mu}(x_i), \boldsymbol{\Psi}^{-1})$ of the underlying latent variables \mathbf{y}_i^* would yield the same class probabilities as the multivariate normal distribution $N_m(a\mathbf{1}_m + c\boldsymbol{\mu}(x_i), c^2\boldsymbol{\Psi}^{-1})$, according to (5.3). Therefore, the multinomial probit model is not identified as there exists more than one set of parameters for which the categorical likelihood $\prod_{i,j} p_{ij}$ is the same.

Identification for the probit model is resolved by setting one restriction on the intercepts $\alpha_1, \dots, \alpha_m$ (location) and $m + 1$ restrictions on the precision matrix $\boldsymbol{\Psi}$ (scale). Restrictions on the intercepts include $\sum_{j=1}^m \alpha_j = 0$ or setting one of the intercepts to zero. In this work, we apply the former restriction to the I-probit model, as this is analogous to the requirement of zero-mean functions in the functional ANOVA decomposition. If A6 holds, then location identification is all that is needed to achieve identification. However, if $\boldsymbol{\Psi}$ is a free parameter to be estimated, only $m(m - 1)/2 - 1$ parameters are identified. Many possible specifications of the restriction on $\boldsymbol{\Psi}$ is possible, depending on the number of alternatives m and the intended effect of $\boldsymbol{\Psi}$, for example:

- **Case $m = 2$** (minimum number of restrictions = 3).

$$\boldsymbol{\Psi} = \begin{pmatrix} 1 & \\ 0 & 0 \end{pmatrix}, \text{ or } \boldsymbol{\Psi} = \begin{pmatrix} 1 & \\ 0 & 1 \end{pmatrix}$$

- **Case $m = 3$** (minimum number of restrictions = 4).

$$\boldsymbol{\Psi} = \begin{pmatrix} 1 & & \\ \psi_{12} & \psi_{22} & \\ 0 & 0 & 0 \end{pmatrix}, \text{ or } \boldsymbol{\Psi} = \begin{pmatrix} 1 & & \\ 0 & \psi_{22} & \\ 0 & 0 & \psi_{33} \end{pmatrix}$$

- **Case $m \geq 4$** (minimum number of restrictions = $m + 1$).

$$\mathbf{\Psi} = \begin{pmatrix} 1 & & & & \\ \psi_{12} & \psi_{22} & & & \\ \vdots & \vdots & \ddots & & \\ \psi_{1,m-1} & \psi_{2,m-1} & \cdots & \psi_{m-1,m-1} & \\ 0 & 0 & \cdots & 0 & 0 \end{pmatrix}, \text{ or } \mathbf{\Psi} = \begin{pmatrix} \psi_{11} & & & & \\ & \psi_{22} & & & \\ & & \ddots & & \\ & & & \psi_{mm} & \end{pmatrix}$$

Remark 5.1. Identification is most commonly achieved by fixing the latent propensities of one of the classes to zero and fixing one element the covariance matrix (Dansie, 1985; Bunch, 1991). Fixing the last class, say, to zero, i.e. $y_{im}^* = 0, \forall i = 1, \dots, n$ has the effect of shrinking $\mathbf{\Psi}$ to $(m - 1) \times (m - 1)$ in size, and thus one more restriction needs to be made (typically, the first element $\mathbf{\Psi}_{11}$ is set to one). This speaks to the fact that the absolute values of the latent propensities themselves do not matter, but their relative differences do—see Section X. We also remark that for the binary case ($m = 2$), setting the latent propensities for the second class to zero and fixing the remaining variance parameter to one yields, for $i = 1, \dots, n$,

$$\begin{aligned} p_{i1} &= P(y_{i1}^* > y_{i2}^* = 0) \\ &= P(\alpha_1 + f_1(x_i) + \epsilon_{i1} > 0 \mid \epsilon_{i1} \stackrel{\text{iid}}{\sim} N(0, 1)) \\ &= \Phi(\alpha_1 + f_1(x_i)) \end{aligned}$$

and $p_{i2} = 1 - \Phi(\alpha_1 + f_1(x_i))$, the familiar binary probit model. Note that in the binary case only one set of latent propensities need to be estimated, so we can drop the subscript ‘1’ in the above equations. In fact, for m classes, only $m - 1$ sets of regression functions need to be estimated (since one of them needs to be fixed), but in the multinomial presentation of this thesis we define regression functions for each class.

Now, we turn to a discussion of the role of $\mathbf{\Psi}$ in the model. In decision theory, the independence axiom states that an agent’s choice between a set of alternatives should not be affected by the introduction or elimination of a choice option. The probit model is suitable for modelling multinomial data where the independence axiom, which is also known as the *independence of irrelevant alternatives* (IIA) assumption, is not desired. Such cases arise frequently in economics and social science, and the famous Red-Bus-Blue-Bus example is often used to illustrate IIA: suppose commuters face the decision between taking cars and red busses. The addition of blue busses to commuters’ choices

should, in theory, be more likely chosen by those who prefer taking the bus over cars. That is, assuming commuters are indifferent about the colour of the bus, commuters who are predisposed to taking the red bus would see the blue bus as an identical alternative. Yet, if IIA is imposed, then the three choices are distinct, and the fact that red and blue busses are substitutable is ignored.

To put it simply, the model is IIA if choice probabilities depend only on the choice in consideration, and not on any other alternatives. In the I-probit model, or rather, in probit models in general, choice dependency is controlled by the error precision matrix Ψ . Specifically, the off-diagonal elements Ψ_{jk} capture the correlation between alternatives j and k . Allowing all $m(m+1)/2$ covariance elements of Ψ to be non-zero leads to the *full I-probit model*, and would not assume an IIA position. Figure 5.1 illustrates the covariance structure for the marginal distribution of the latent propensities, $\mathbf{V}_{y^*} = \Psi \otimes \mathbf{H}_\eta^2 + \Psi^{-1} \otimes \mathbf{I}_n$, and of the I-prior $\mathbf{V}_f = \Psi \otimes \mathbf{H}_\eta^2$.

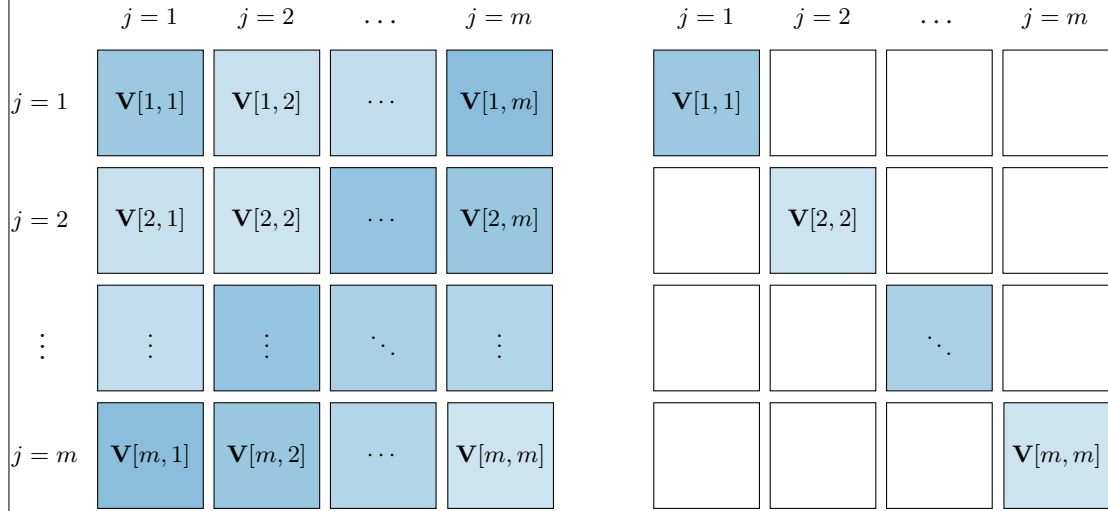


Figure 5.1: Illustration of the covariance structure of the full I-probit model (left) and the independent I-probit model (right). The full model has m^2 blocks of $n \times n$ symmetric matrices, and the blocks themselves are arranged symmetrically about the diagonal. The independent model, on the other hand, has a block diagonal structure, and its sparsity induces simpler computational methods for estimation.

fig:iprobco
vstr

While it is an advantage to be able to model the correlations across choices (unlike in logistic models), there are applications where the IIA assumption would not adversely affect the analysis, such as classification tasks. Some analyses might also be indifferent

as to whether or not choice dependency exists. In these situations, it would be beneficial, algorithmically speaking, to reduce the I-probit model to a simpler version by assuming $\Psi = \text{diag}(\psi_1, \dots, \psi_m)$, which would trigger the IIA assumption in the I-probit model. We refer to this model as the *independent I-probit model*. The independence assumption causes the distribution of the latent variables to be $y_{ij}^* \sim N(\mu_k(x_i), \sigma_j^2)$ for $j = 1, \dots, m$, where $\sigma_j^2 = \psi_j^{-1}$. As a continuation of line (5.3), we can show the class probability p_{ij} to be

$$\begin{aligned} p_{ij} &= \int \cdots \int \prod_{\substack{k=1 \\ \{y_{ij}^* > y_{ik}^* | \forall k \neq j\}}}^m \left\{ \phi(y_{ik}^* | \mu_k(x_i), \sigma_k^2) dy_{ik}^* \right\} \\ &= \int \prod_{\substack{k=1 \\ k \neq j}}^m \Phi\left(\frac{y_{ij}^* - \mu_k(x_i)}{\sigma_k}\right) \cdot \phi(y_{ij}^* | \mu_j(x_i), \sigma_j^2) dy_{ij}^* \\ &= E_Z \left[\prod_{\substack{k=1 \\ k \neq j}}^m \Phi\left(\frac{\sigma_j}{\sigma_k} Z + \frac{\mu_j(x_i) - \mu_k(x_i)}{\sigma_k}\right) \right] \end{aligned} \quad (5.5)$$

{eq:pij2}

where $Z \sim N(0, 1)$, $\Phi(\cdot)$ its cdf, and $\phi(\cdot | \mu, \sigma^2)$ is the pdf of $X \sim N(\mu, \sigma^2)$. The equation (5.3) is thus simplified to a unidimensional integral involving the Gaussian pdf and cdf, which can be computed fairly efficiently using quadrature methods. The probit link function is evidently seen in the above equation.

5.3 Estimation

As with the normal I-prior model, an estimate of the posterior regression function with optimised hyperparameters is sought. The log likelihood function $L(\cdot)$ for θ using all n observations $\{(y_1, x_1), \dots, (y_n, x_n)\}$ is obtained by performing the following integration:

$$L(\theta) = \log \iint p(\mathbf{y} | \mathbf{y}^*, \theta) p(\mathbf{y}^* | \mathbf{w}, \theta) p(\mathbf{w} | \theta) d\mathbf{y}^* d\mathbf{w} \quad (5.6)$$

{eq:iprobitt
lik}

Here, $p(\mathbf{w} | \theta)$ is the pdf of $MN_{n,m}(\mathbf{0}, \mathbf{I}_n, \Psi)$, $p(\mathbf{y}^* | \mathbf{w}, \theta)$ is the pdf of $MN_{n,m}(\mathbf{1}_n \boldsymbol{\alpha}^\top + \mathbf{H}_\eta \mathbf{w}, \mathbf{I}_n, \Psi^{-1})$, and $p(\mathbf{y} | \mathbf{y}^*, \theta) = \prod_{i=1}^n \prod_{j=1}^m [y_{ij}^* = \max \mathbf{y}_{i\cdot}^*]^{[y_i=j]}$, with $0^0 := 1$. Note that, given the corresponding latent propensities $\mathbf{y}_{i\cdot}^* = (y_{i1}^*, \dots, y_{im}^*)^\top$, the distribution $y_i | \mathbf{y}_{i\cdot}^*$ is tantamount to a degenerate categorical distribution, since after knowing which of the latent propensities is largest, knowledge of the outcome of the categorical response becomes a certainty.

Marginalising out the random effects \mathbf{w} from the pdf of the latent propensities is a relatively simple task as this involves two normal densities, but simplicity ends there. Continuing on (5.6), we get

$$\begin{aligned} L(\theta) &= \log \int p(\mathbf{y}|\mathbf{y}^*, \theta) p(\mathbf{y}^*|\theta) d\mathbf{y}^* \\ &= \log \int \prod_{i=1}^n [y_{iy_i}^* = \max \mathbf{y}_i^*] \phi(\mathbf{y}^* | \mathbf{1}_n \boldsymbol{\alpha}^\top, \boldsymbol{\Psi} \otimes \mathbf{H}_\eta^2 + \boldsymbol{\Psi}^{-1} \otimes \mathbf{I}_n) d\mathbf{y}^*. \end{aligned}$$

Being an integral of order nm , it is an incredibly daunting problem to tackle which hampers hope of obtaining parameter estimates via direct likelihood maximisation methods. Furthermore, the posterior distribution of the regression function, which requires the density of $\mathbf{w}|\mathbf{y}$, obtained via

$$\begin{aligned} p(\mathbf{w}|\mathbf{y}) &\propto p(\mathbf{y}|\mathbf{w})p(\mathbf{w}) \\ &= \int p(\mathbf{y}|\mathbf{y}^*, \mathbf{w})p(\mathbf{y}^*|\mathbf{w})p(\mathbf{w})d\mathbf{y}^*, \end{aligned}$$

has normalising constant equal to the marginalisation integral in (5.6). The challenge of estimation is then to first overcome this intractability by means of a suitable approximation of the integral. We present three possible avenues to achieve this aim, namely the Laplace approximation, variational Bayes, and Markov chain Monte Carlo (MCMC) methods.

5.3.1 Laplace approximation

An alternative approach to compute the log-likelihood function is to bypass the latent propensities completely, i.e.

$$\begin{aligned} L(\theta) &= \log \int p(\mathbf{y}|\mathbf{w}, \theta) p(\mathbf{w}|\theta) d\mathbf{w} \\ &= \log \int \prod_{i=1}^n \prod_{j=1}^m \left(g_j^{-1}(\alpha_k + \overbrace{f_k(x_i)}^{\sum_{i'=1}^n h_\eta(x_i, x_{i'}) w_{i'k}})_{k=1}^m \right)^{[y_i=j]} \cdot \text{MN}_{n,m}(\mathbf{w}|\mathbf{0}, \mathbf{I}_n, \boldsymbol{\Psi}) d\mathbf{w} \quad (5.7) \end{aligned}$$

{eq:intractablelikelihood}

where we have denoted the class probabilities p_{ij} from (5.3) using the function $g_j^{-1} : \mathbb{R}^m \rightarrow [0, 1]$, which highlights the dependence of the class probabilities on \mathbf{w} . This formulation sets us up nicely to apply Laplace's method of approximating the integral in (5.7).

The focus here is to obtain the posterior density $p(\mathbf{w}|\mathbf{y}) \propto p(\mathbf{y}|\mathbf{w})p(\mathbf{w}) =: e^{R(\mathbf{w})}$ with normalising constant equal to the marginal density of \mathbf{y} , $p(\mathbf{y}) = \int e^{R(\mathbf{w})} d\mathbf{w}$. Laplace's method (Kass and Raftery, 1995, §4.1.1, pp. 777–778) entails expanding a Taylor series for R about its posterior mode $\hat{\mathbf{w}} = \arg \max_{\mathbf{w}} p(\mathbf{y}|\mathbf{w})p(\mathbf{w})$, which gives the relationship

$$\begin{aligned} R(\mathbf{w}) &= R(\hat{\mathbf{w}}) + \underbrace{(\mathbf{w} - \hat{\mathbf{w}})^\top \nabla R(\hat{\mathbf{w}})}_{\rightarrow 0} - \frac{1}{2}(\mathbf{w} - \hat{\mathbf{w}})^\top \boldsymbol{\Omega}(\mathbf{w} - \hat{\mathbf{w}}) + \dots \\ &\approx R(\hat{\mathbf{w}}) + -\frac{1}{2}(\mathbf{w} - \hat{\mathbf{w}})^\top \boldsymbol{\Omega}(\mathbf{w} - \hat{\mathbf{w}}), \end{aligned}$$

because, assuming that R has a unique maximum, ∇R evaluated at its mode is zero. This is recognised as the logarithm of an unnormalised Gaussian density, implying $\mathbf{w}|\mathbf{y} \sim N_n(\hat{\mathbf{w}}, \boldsymbol{\Omega}^{-1})$. Here, $\boldsymbol{\Omega} = -\nabla^2 R(\mathbf{w})|_{\mathbf{w}=\hat{\mathbf{w}}}$ is the negative Hessian of Q evaluated at the posterior mode, and is typically obtained as a byproduct of the maximisation routine of R using gradient or quasi-gradient based methods.

The marginal distribution is then approximated by

$$\begin{aligned} p(\mathbf{y}) &= \int \exp \underbrace{R(\mathbf{w})}_{\approx R(\hat{\mathbf{w}}) - \frac{1}{2}(\mathbf{w} - \hat{\mathbf{w}})^\top \boldsymbol{\Omega}(\mathbf{w} - \hat{\mathbf{w}})} d\mathbf{w} \\ &\approx (2\pi)^{n/2} |\boldsymbol{\Omega}|^{-1/2} e^{Q(\hat{\mathbf{w}})} \int (2\pi)^{-n/2} |\boldsymbol{\Omega}|^{1/2} \exp \left(-\frac{1}{2}(\mathbf{w} - \hat{\mathbf{w}})^\top \boldsymbol{\Omega}(\mathbf{w} - \hat{\mathbf{w}}) \right) d\mathbf{w} \\ &= (2\pi)^{n/2} |\boldsymbol{\Omega}|^{-1/2} p(\mathbf{y}|\hat{\mathbf{w}})p(\hat{\mathbf{w}}). \end{aligned}$$

The log marginal density of course depends on the parameters θ , which becomes the objective function to maximise in a likelihood maximising approach. Note that, should a fully Bayesian approach be undertaken, i.e. priors prescribed on the model parameters using $\theta \sim p(\theta)$, then this approach is viewed as a maximum a posteriori approach.

In any case, each evaluation of the objective function $L(\theta) = \log p(\mathbf{y}|\theta)$ involves finding the posterior modes $\hat{\mathbf{w}}$. This is a slow and difficult undertaking, especially for large sample sizes n —even assuming computation of the class probabilities g^{-1} is efficient—because the dimension of this integral is exactly the sample size. Furthermore, obtaining standard errors for the parameters are cumbersome, and it is likely that a computationally burdensome bootstrapping approach is needed. Lastly, as a comment, Laplace's method only approximates the true marginal likelihood well if the true function is small far away from the mode.

5.3.2 Variational EM algorithm

We turn to variational methods as a means of approximating the posterior densities of interest and obtain parameter estimates. Variational methods are widely discussed in the machine learning literature, but there have been efforts to popularise it in statistics (Blei et al., 2017). Although variational inference is typically seen as a fully Bayesian method, whereby approximate posterior densities are sought for the latent variables and parameters, our goal is to apply variational inference to facilitate a pseudo maximum likelihood approach.

Consider employing an EM algorithm, similar to the one seen in the previous chapter, to estimate I-probit models. This time, treat both the latent propensities \mathbf{y}^* and the I-prior random effects \mathbf{w} as ‘missing’, so the complete data is $\{\mathbf{y}, \mathbf{y}^*, \mathbf{w}\}$. Now, due to the independence of the observations $i = 1, \dots, n$, the complete data log-likelihood is

$$\begin{aligned} \log p(\mathbf{y}, \mathbf{y}^*, \mathbf{w}) &= \sum_{i=1}^n \left\{ \log p(y_i | \mathbf{y}_{i\cdot}^*) + \log p(\mathbf{y}_{i\cdot}^* | \mathbf{w}_{i\cdot}) + \log p(\mathbf{w}_{i\cdot}) \right\} \\ &= -\frac{1}{2} \sum_{i=1}^n [y_{ij}^* = \max_k y_{ik}^*] \left\{ (\mathbf{y}_{i\cdot}^* - \boldsymbol{\alpha} - \mathbf{w}_{i\cdot}^\top \mathbf{h}_\eta(x_i))^\top \boldsymbol{\Psi} (\mathbf{y}_{i\cdot}^* - \boldsymbol{\alpha} - \mathbf{w}_{i\cdot}^\top \mathbf{h}_\eta(x_i)) \right. \\ &\quad \left. + \mathbf{w}_{i\cdot}^\top \boldsymbol{\Psi}^{-1} \mathbf{w}_{i\cdot} \right\} + \text{const.} \end{aligned}$$

which looks like the complete data log-likelihood seen previously in (4.9), except that here, together with the $\mathbf{w}_{i\cdot}$ ’s, the $\mathbf{y}_{i\cdot}^*$ ’s are never observed.

For the E-step, it is of interest to determine the posterior density $p(\mathbf{y}^*, \mathbf{w} | \mathbf{y}) = p(\mathbf{y}^* | \mathbf{w}, \mathbf{y}) p(\mathbf{w} | \mathbf{y})$, which we have discerned from the discussion at the beginning of this section that this is hard to obtain, since it involves an intractable marginalising integral. We thus seek a suitable approximation

$$p(\mathbf{y}^*, \mathbf{w} | \mathbf{y}, \theta) \approx \tilde{q}(\mathbf{y}^*, \mathbf{w}),$$

where \tilde{q} satisfies $\tilde{q} = \arg \min_q \text{KL}(q || p)$, subject to certain constraints. The constraint considered by us in this thesis is that q satisfies a *mean-field* factorisation

$$q(\mathbf{y}^*, \mathbf{w}, \theta) = q(\mathbf{y}^*) q(\mathbf{w}).$$

Under this scheme, the variational distribution for \mathbf{y}^* is found to be a *conically truncated multivariate normal* distribution, and for \mathbf{w} , a multivariate normal distribution.

It can be shown that, for some variational density q , the marginal log-likelihood is an upper-bound for the quantity $\mathcal{L}_\theta(q)$

$$\log p(\mathbf{y}|\theta) \geq \mathbb{E}_{\mathbf{y}^*, \mathbf{w} \sim q} \log p(\mathbf{y}, \mathbf{y}^*, \mathbf{w}|\theta) - \mathbb{E}_{\mathbf{y}^*, \mathbf{w} \sim q} \log q(\mathbf{y}^*, \mathbf{w}) =: \mathcal{L}_\theta(q),$$

a quantity often referred to as the *evidence lower bound* (ELBO). It turns out that minimising $\text{KL}(q||p)$ is equivalent to maximising the ELBO, a quantity that is more practical to work with than the KL divergence, and certainly more tractable than the log marginal density. Hence, if q approximates the true posterior well, then the ELBO is a suitable proxy for the marginal log-likelihood.

In practice, obtaining parameter estimates which maximise the ELBO and the approximate posterior distribution $q(\mathbf{y}^*, \mathbf{w})$ is achieved using a variational EM algorithm, an EM algorithm in which the conditional distributions are replaced with a variational approximation. The t 'th E-step entails obtaining the density $q^{(t+1)}$ as a solution to $\arg \max_q \mathcal{L}_\theta(q)$, keeping θ fixed at the current estimate $\theta^{(t)}$. Let $\bar{\mathbf{y}}^* = \mathbf{y}^* - \mathbf{1}_n \boldsymbol{\alpha}^\top$. The objective function to be maximised is computed as

$$\begin{aligned} Q(\theta) &= \mathbb{E}_{\mathbf{y}^*, \mathbf{w} \sim q^{(t+1)}} [\log p(\mathbf{y}, \mathbf{y}^*, \mathbf{w}|\theta)] \\ &= \text{const.} - \frac{1}{2} \text{tr} \mathbb{E}_{\mathbf{y}^*, \mathbf{w} \sim q} \left[\boldsymbol{\Psi} (\bar{\mathbf{y}}^{*\top} \bar{\mathbf{y}}^* + \mathbf{w}^\top \mathbf{H}_\eta^2 \mathbf{w} - 2\mathbf{w}^\top \mathbf{H}_\eta \bar{\mathbf{y}}^*) + \boldsymbol{\Psi}^{-1} \mathbf{w}^\top \mathbf{w} \right], \end{aligned}$$

and this is maximised with respect to θ in the M-step to obtain $\theta^{(t+1)}$, and the process repeated until convergence of the ELBO. A full derivation of the variational EM algorithm used by us will be described in [Section 5.4](#).

5.3.3 Markov chain Monte Carlo methods

As an alternative to the deterministic Bayesian approach of variational inference, it is possible to use Markov chain Monte Carlo sampling methods as an approach to stochastically approximate the intractable posterior distribution.

[Albert and Chib \(1993\)](#) showed that the latent variable approach to probit models can be analysed using exact Bayesian methods, due to the underlying normality structure. Paired with corresponding conjugate prior choices, sampling from the posterior is

very simple using a Gibbs sampling approach. On the other hand, this data augmentation scheme enlarges the variable space to $n + q$ dimensions, where q is the number of parameters to estimate, which is inefficient and computationally challenging especially when n is large. It is no longer possible to marginalise the normal latent variables from the model, as this is intractable, as discussed previously.

Hamiltonian Monte Carlo is another possibility, since it does not require conjugacy. For binary models, this is a feasible approach because the class probabilities are a function of the normal CDF, which means that it is doable using off-the-shelf software such as Stan. However, with multinomial responses, the arduous task of computing class probabilities, which involve integration of an at most m -dimensional normal density, must be addressed separately.

5.3.4 Comparison of estimation methods

In this subsection, we utilise a toy binary classification data set which has been simulated according to a spiral pattern, as in Figure 5.2. The predictor variables are X_1 and X_2 , each of which are scaled similarly.




Figure 5.2: A plot of simulated spiral data set.

fig:example
iprobit

The I-probit model that is fitted is

$$\begin{aligned} y_i &\sim \text{Bern}(p_i) \\ \Phi^{-1}(p_i) &= \alpha + \sum_{k=1}^n h_{\lambda}(x_i, x_k) w_k \\ w_1, \dots, w_n &\stackrel{\text{iid}}{\sim} \text{N}(0, 1). \end{aligned}$$

This binary model follows from the more general multinomial I-probit model by fixing all latent propensities in one of the classes to zero, and setting $\Psi = \mathbf{I}_m$. This is possible because only differences in latent propensities are of interest, and not the actual values themselves, and thus only $m - 1$ sets of posterior regression functions need to be estimated—see ?? for further details.

The three estimation methods described aim to overcome the intractable integral by means of either a deterministic approximation (Laplace and variational inference) or a stochastic approximation (Hamiltonian MC). For the Bayesian methods, i.e. variational inference and Hamiltonian MC, vague priors were used on α and λ , namely $\text{N}(0, 100)$ and $\text{N}_+(0, 100)$ respectively. Restriction of λ to the positive orthant is required for identifiability. The Laplace and variational methods were performed in the **iprobbit** package, while **Stan** was used to code the Hamiltonian MC sampler. The results are presented in Table 5.1.

Table 5.1: Table comparing the estimated parameter values, (marginal) log-likelihood values, and also time taken for the three estimation methods.

	Laplace approximation	Variational inference	Hamiltonian MC
Intercept (α)	-0.02 (0.03)	0.00 (0.06)	0.00 (0.58)
Scale (λ)	0.85 (0.01)	5.67 (0.23)	29.3 (5.21)
Log density	-202.7	-140.7	-163.8
Error rate (%)	44.7	0.00	2.24
Brier score	0.20	0.02	0.01
Iterations	20	56	2000
Time taken (s)	>3600	5.32	>3600

The three methods pretty much concur on the estimation of the intercept, but not on the RKHS scale parameter. As a result, the log-density value at the optima is also different in all three methods. Notice the high posterior standard deviation for the scale

tab:comprei
probit

parameter in the HMC method. The posterior density for λ was very positively skewed, and this contributed to the large posterior mean.

A plot of the log-likelihood surface for three methods in [Figure 5.3](#) reveals some insight. The variational likelihood reveals two ridges, with the maxima occurring around the intersection of these two ridges. The Laplace likelihood seems to indicate a similar shape—in both the Laplace and variational method, the posterior distribution of \mathbf{w} is approximated by a Gaussian distribution, with different means and variances. However, parts of the Laplace likelihood are poorly approximated resulting in a loss of fidelity around the supposed maxima, which might have contributed to the set of values that were estimated. Laplace’s method is known to yield poor approximations to probit-type likelihoods, as studied by [Kuss and Rasmussen \(2005\)](#). On the other hand, the log-likelihood using the posterior distribution of the Hamiltonian MC sampler (treating parameters as fixed values) yields a completely different shape compared to the other two methods.

In terms of predictive abilities, both the variational and Hamiltonian MC methods, even though the posteriors are differently estimated, has good predictive performance as indicated by their error rates and Brier scores. [Figure 5.3](#) shows that HMC is more confident of new data predictions compared to variational inference, as indicated by the intensity of the shaded regions (HMC is stronger than VI). Laplace’s method gave poor predictive performance.

Finally, on the computational side, variational inference was by far the fastest method to fit the model. Sampling using Hamiltonian MC was very slow, because the parameter space is in effect $O(n + 2)$ (parameters are $\{w_1, \dots, w_n, \alpha, \lambda\}$), and unlike in the normal model, we are not able to easily marginalise out the I-prior. As for Laplace, each Newton step involves obtaining posterior modes of the w ’s, and this contributed to the slowness of this method. The reality is that variational inference takes seconds to complete what either the Laplace or full MCMC methods would take hours to. The predictive performance, while not as good as HMC, is certainly an acceptable compromise in favour of speed.

5.4 A variational algorithm

We present a variational inference algorithm to estimate the I-probit latent variables \mathbf{y}^* and \mathbf{w} , together with the parameters $\theta = \{\boldsymbol{\alpha} = (\alpha_1, \dots, \alpha_m)^\top, \eta\}$ with *fixed error*

sec:iprobit
var

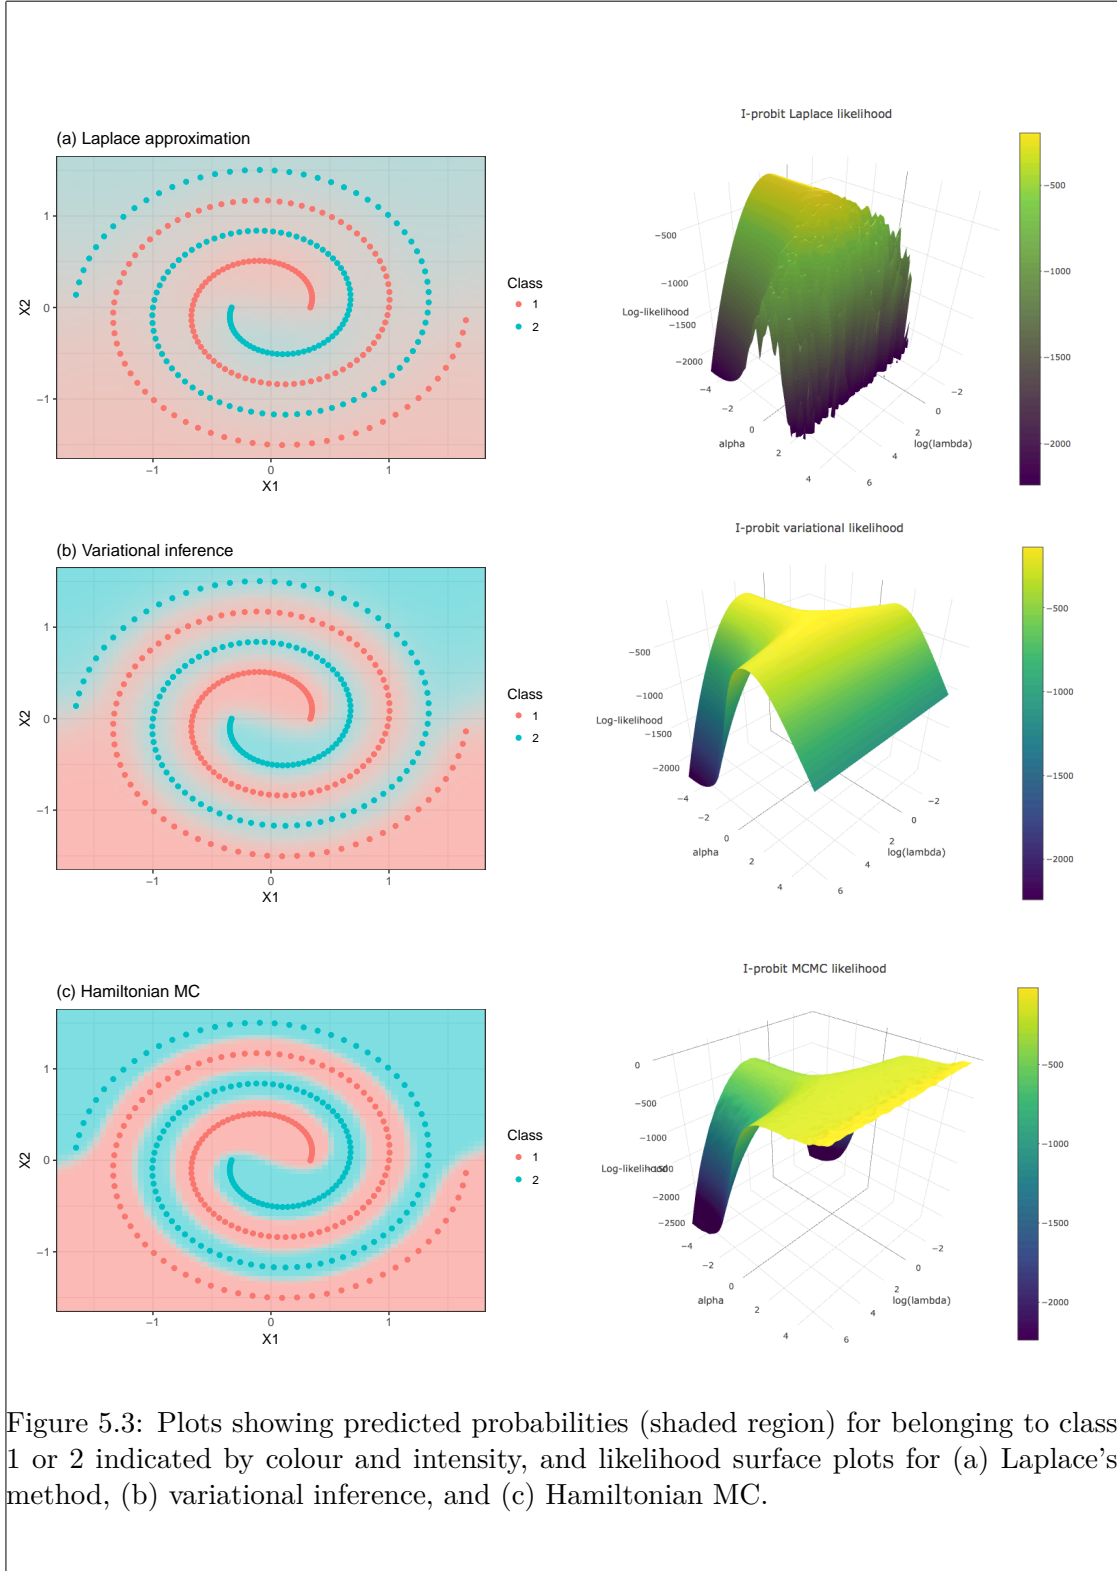


fig:example
iprobitfit

precision Ψ^2 . Begin by choosing prior distributions on the parameters, $p(\theta) = p(\boldsymbol{\alpha})p(\eta)$. The following flat, uninformative priors are suggested:

- **Kernel parameters** η . This may include parameters such as the Hurst index, lengthscale and offset parameters, in addition to the RKHS scale parameters $\lambda_1, \dots, \lambda_p$, and each with their own support. For the scale parameters, assign each λ_k the vague prior

$$\lambda_k \stackrel{\text{iid}}{\sim} \text{N}(0, v_\lambda = 0.001^{-1}), \quad k = 1, \dots, p.$$

As $v_k^{-1} \rightarrow 0$, the prior becomes $p(\lambda_k) \propto \text{const.}$, an improper prior. The default choice for the rest of the kernel parameters is an improper prior $p(\eta) \propto \text{const.}$

- **Intercepts** $\alpha_1, \dots, \alpha_m$. Assign independent, vague normal priors for each intercept

$$\alpha_j \stackrel{\text{iid}}{\sim} \text{N}(0, v_\alpha = 0.001^{-1}).$$

Although one may devote more attention to the prior specification of these parameters, for our purposes it suffices that they are independent component-wise, and that they are conjugate priors for the complete conditional density $p(\theta|\mathbf{y}, \mathbf{y}^*, \mathbf{w})$.

The posterior density of $\mathcal{Z} = \{\mathbf{y}^*, \mathbf{w}, \theta\}$ is approximated by a mean-field variational density q , i.e.

$$p(\mathbf{y}^*, \mathbf{w}, \theta|\mathbf{y}) = q(\mathbf{y}^*)q(\mathbf{w})q(\theta).$$

Additionally, we assume independence among the components of θ so that $q(\theta) = \prod_k q(\theta_k)$. We now present the mean-field variational distributions for each of unknowns in \mathcal{Z} . On notation: we will typically refer to posterior means of the parameters \mathbf{y}^* , \mathbf{w} , θ and so on by the use of a tilde. For instance, we write $\tilde{\mathbf{w}}$ to mean $\text{E}_{\mathbf{w} \sim q}[\mathbf{w}]$, the expected value of \mathbf{w} under the pdf $q(\mathbf{w})$. The distributions are simply stated, but a full derivation is given in the appendix.

²It turns out that the variational algorithm as presented is not suited to estimate Ψ . This issue is discussed further in Section X.




Figure 5.4: A DAG of the I-probit model. Observed/fixed nodes are shaded, while double-lined nodes represents calculable quantities.

5.4.1 Latent propensities \mathbf{y}^*

The fact that the rows $\mathbf{y}_i^* \in \mathbb{R}^m$, $i = 1, \dots, n$ of $\mathbf{y}^* \in \mathbb{R}^{n \times m}$ are independent can be exploited, which yields an induced factorisation $q(\mathbf{y}^*) = \prod_{i=1}^n q(\mathbf{y}_i^*)$. Define the set $\mathcal{C}_j = \{y_{ij}^* > y_{ik}^* \mid \forall k \neq j\}$. Then $q(\mathbf{y}_i^*)$ is the density of a multivariate normal distribution with mean $\tilde{\boldsymbol{\mu}}_{i.} = \tilde{\boldsymbol{\alpha}} + \tilde{\mathbf{w}}^\top \tilde{\mathbf{h}}_\eta(x_i)$, and variance $\boldsymbol{\Psi}^{-1}$ subject to the truncation of its components to the set \mathcal{C}_{y_i} . That is, for each $i = 1, \dots, n$ and noting the observed value $y_i \in \{1, \dots, m\}$, the \mathbf{y}_i^* 's are distributed according to

$$\mathbf{y}_i^* \stackrel{\text{iid}}{\sim} \begin{cases} N_m(\tilde{\boldsymbol{\mu}}_{i.}, \boldsymbol{\Psi}^{-1}) & \text{if } y_{iy_i}^* > y_{ik}^*, \forall k \neq y_i \\ 0 & \text{otherwise.} \end{cases} \quad (5.8)$$

{eq:ystandardi
st}

We denote this by $\mathbf{y}_i^* \stackrel{\text{iid}}{\sim} \text{tN}(\tilde{\boldsymbol{\mu}}_{i.}, \boldsymbol{\Psi}^{-1}, \mathcal{C}_{y_i})$, and the important properties of this distribution are explored in the appendix.

The required expectations $E\mathbf{y}_i^* = E(y_{i1}^*, \dots, y_{im}^*)^\top$ are tricky to compute. One strategy might be Monte Carlo integration: using samples from $N_m(\tilde{\boldsymbol{\mu}}_{i.}, \boldsymbol{\Psi}^{-1})$, disregard those that do not satisfy the condition $y_{iy_i}^* > y_{ik}^*, \forall k \neq y_i$, and then take the sample average. This works reasonably well so long as the truncation region does not fall into the extreme tails of the multivariate normal. Alternatively, a fast, Gibbs based approach

to estimating the mean or any other quantity $\mathbb{E}[r(\mathbf{y}_{i.}^*)]$ can be implemented, and this is detailed in the appendix.

If the independent I-probit model is considered, where the covariance matrix has the independent structure $\mathbf{\Psi} = \text{diag}(\sigma_1^{-2}, \dots, \sigma_m^{-2})$, then the expected value can be considered component-wise, and each component of this expectation is given by

$$\tilde{y}_{ik}^* = \begin{cases} \tilde{\mu}_{ik} - \sigma_k C_i^{-1} \int \phi_{ik}(z) \prod_{l \neq k, j} \Phi_{il}(z) \phi(z) \, dz & \text{if } k \neq y_i \\ \tilde{\mu}_{iy_i} - \sigma_{y_i} \sum_{k \neq y_i} (\tilde{y}_{ik}^* - \tilde{\mu}_{ik}) & \text{if } k = y_i \end{cases} \quad (5.9)$$

{eq:ystart
date}

with

$$\begin{aligned} \phi_{ik}(Z) &= \phi\left(\frac{\sigma_{y_i}}{\sigma_k} Z + \frac{\tilde{\mu}_{iy_i} - \tilde{\mu}_{ik}}{\sigma_k}\right) \\ \Phi_{ik}(Z) &= \Phi\left(\frac{\sigma_{y_i}}{\sigma_k} Z + \frac{\tilde{\mu}_{iy_i} - \tilde{\mu}_{ik}}{\sigma_k}\right) \\ C_i &= \int \prod_{l \neq j} \Phi_{il}(z) \phi(z) \, dz \end{aligned}$$

and $Z \sim \mathcal{N}(0, 1)$ with pdf and cdf $\phi(\cdot)$ and $\Phi(\cdot)$ respectively. The integrals that appear above are functions of a unidimensional Gaussian pdf, and these can be computed rather efficiently using quadrature methods.

5.4.2 I-prior random effects w

Given that both $\text{vec } \mathbf{y}^* | \text{vec } \mathbf{w}$ and $\text{vec } \mathbf{w}$ are normally distributed, we find that the conditional posterior distribution $p(\mathbf{w} | \mathcal{Z}_{-\mathbf{w}}, \mathbf{y})$ is also normal, and therefore the approximate posterior density q for $\text{vec } \mathbf{w} \in \mathbb{R}^{nm}$ is also normal with mean and precision given by

$$\text{vec } \tilde{\mathbf{w}} = \tilde{\mathbf{V}}_w (\mathbf{\Psi} \otimes \tilde{\mathbf{H}}_\eta) \text{vec}(\tilde{\mathbf{y}}^* - \mathbf{1}_n \tilde{\boldsymbol{\alpha}}^\top) \quad \text{and} \quad \tilde{\mathbf{V}}_w^{-1} = (\mathbf{\Psi} \otimes \tilde{\mathbf{H}}_\eta^2) + (\mathbf{\Psi}^{-1} \otimes \mathbf{I}_n). \quad (5.10)$$

{eq:varipos
tw}

We note the similarity between (5.10) above and the posterior distribution for the I-prior random effects in a normal model (4.7) seen in the previous chapter. Naïvely computing the inverse $\tilde{\mathbf{V}}_w^{-1}$ presents a computational challenge, as this takes $O(n^3 m^3)$ time. By exploiting the Kronecker product structure in $\tilde{\mathbf{V}}_w$, we are able to efficiently compute the required inverse in roughly $O(n^3 m)$ time—see the appendix for details.

If the independent I-probit model is assumed, i.e. $\tilde{\Psi} = \text{diag}(\tilde{\psi}_1, \dots, \tilde{\psi}_m)$, then the posterior covariance matrix $\tilde{\mathbf{V}}_w$ has a simpler structure: random matrix \mathbf{w} will have columns which are independent of each other. By writing $\mathbf{w}_{\cdot j} = (w_{1j}, \dots, w_{nj})^\top \in \mathbb{R}^n$, $j = 1, \dots, m$, to denote the column vectors of \mathbf{w} and with a slight abuse of notation, we have that

$$N_{nm}(\text{vec } \mathbf{w} | \text{vec } \tilde{\mathbf{w}}, \tilde{\mathbf{V}}_w) = \prod_{j=1}^m N_n(\mathbf{w}_{\cdot j} | \tilde{\mathbf{w}}_{\cdot j}, \tilde{\mathbf{V}}_{w_j}),$$

where

$$\tilde{\mathbf{w}}_{\cdot j} = \psi_j \tilde{\mathbf{V}}_{w_j} \tilde{\mathbf{H}}_\eta (\tilde{\mathbf{y}}_j^* - \tilde{\alpha}_j \mathbf{1}_n) \quad \text{and} \quad \tilde{\mathbf{V}}_{w_j} = (\psi_j \tilde{\mathbf{H}}_\eta^2 + \psi_j^{-1} \mathbf{I}_n)^{-1}.$$

The consequence of this is that the posterior regression functions are class independent, the exact intended effect by specifying a diagonal precision matrix Ψ .

5.4.3 Kernel parameters η

The posterior density q involving the kernel parameters is of the form

$$\log q(\eta) = -\frac{1}{2} \text{tr} \mathbb{E}_{\mathcal{Z} \setminus \{\eta\} \sim q} \left[(\mathbf{y}^* - \mathbf{1}_n \boldsymbol{\alpha}^\top - \mathbf{H}_\eta \mathbf{w}) \Psi (\mathbf{y}^* - \mathbf{1}_n \boldsymbol{\alpha}^\top - \mathbf{H}_\eta \mathbf{w})^\top \right] + \log p(\eta) + \text{const.}$$

where $p(\eta)$ is an appropriate prior density for η . Generally, samples $\eta^{(1)}, \dots, \eta^{(T)}$ from $\tilde{q}(\eta)$ may be obtained using a Metropolis algorithm, so that quantities such as $\tilde{\mathbf{H}}_\eta = \mathbb{E}_{\eta \sim q} \mathbf{H}_\eta$ and the like may be approximated using $\frac{1}{T} \sum_{t=1}^T \mathbf{H}_{\eta^{(t)}}$. Details of the Metropolis sampler is available in the appendix.

When only RKHS scale parameters are involved, then the distribution q can be found in closed-form, much like in the exponential family EM algorithm described in [Section 4.3.3](#). Under the same setting as in that subsection, assume that only $\eta = \{\lambda_1, \dots, \lambda_p\}$ need be estimated, and for each $k = 1, \dots, p$, we can decompose the kernel matrix as $\mathbf{H}_\eta = \lambda_k \mathbf{R}_k + \mathbf{S}_k$ and its square as $\mathbf{H}_\eta^2 = \lambda_k^2 \mathbf{R}_k^2 + \lambda_k \mathbf{U}_k + \mathbf{S}_k^2$. Additionally, we impose a further mean-field restriction on $q(\eta)$, i.e., $q(\eta) = \prod_{k=1}^p p(\lambda_k)$. Then, by using independent and identical normal priors on the λ_k 's, such as the one listed at the beginning of this section, we find that $q(\lambda_k)$ is the density of a normal distribution with

sec:varupdate
ta

mean $d_k c_k^{-1}$ and variance c_k^{-1} , where

$$c_k = \text{tr}(\Psi \mathbf{E}[\mathbf{w}^\top \mathbf{R}_k^2 \mathbf{w}]) + v_\lambda^{-2}$$

and

$$d_k = \text{tr}\left(\Psi(\tilde{\mathbf{y}}^* - \mathbf{1}_n \tilde{\boldsymbol{\alpha}}^\top)^\top \mathbf{R}_k \tilde{\mathbf{w}} - \frac{1}{2} \Psi \mathbf{E}[\mathbf{w}^\top \mathbf{U}_k \mathbf{w}]\right).$$

For a method of evaluating quantities such as $\text{tr}(\mathbf{C} \mathbf{E}[\mathbf{w}^\top \mathbf{D} \mathbf{w}])$ for suitably sized matrices \mathbf{C} and \mathbf{D} , refer to the appendix.

5.4.4 Intercepts $\boldsymbol{\alpha}$

Finally, the posterior distribution for the intercepts follow a normal distribution with the normal priors specified earlier. The posterior mean and variance for the intercepts are given by $\tilde{\boldsymbol{\alpha}} = \tilde{\mathbf{A}}^{-1} \tilde{\mathbf{a}}$ and $\tilde{\mathbf{A}}^{-1}$ respectively, where

$$\tilde{\mathbf{a}} = \sum_{i=1}^n \Psi(\tilde{\mathbf{y}}_{i\cdot}^* - \tilde{\mathbf{w}}^\top \tilde{\mathbf{h}}_\eta(x_i)) \quad \text{and} \quad \tilde{\mathbf{A}} = n\Psi + v_\alpha \mathbf{I}_m.$$

If Ψ is diagonal, the components of $\boldsymbol{\alpha}$ would be independent, and each would be distributed according to

$$\text{N}\left(\frac{\psi_j \sum_{i=1}^n (\tilde{y}_{ij}^* - \tilde{f}_{ij})^2}{n\psi_j + v_\alpha^{-1}}, \frac{1}{n\psi_j + v_\alpha^{-1}}\right).$$

Here, we used the notation \tilde{f}_{ij} to mean the (i, j) 'th element of $\mathbf{E}[\mathbf{H}_\eta \mathbf{w}] = \tilde{\mathbf{H}}_\eta \tilde{\mathbf{w}} \in \mathbb{R}^{n \times m}$. Note that it is necessary, as discussed earlier, that $\sum_{j=1}^m \alpha_j = 0$ for identifiability.

5.4.5 The CAVI algorithm

One will have noticed that the evaluation of each component of the posterior depends on knowing the posterior distribution of the rest of the components. This circular dependence is dealt with by way of an iterative updating scheme of the components. Using an arbitrary starting value, each component is updated in turn according to the above derivations, until a maximum number of iterations is reached, or ideally, until a convergence criterion is met. In variational inference, the ELBO is used to assess convergence.

The expression for the ELBO for the I-probit model is derived in the appendix. The CAVI algorithm for the I-probit model is summarised in [Algorithm 1](#).

Similar to the EM algorithm, each iteration of the algorithm increases the ELBO to a stationary point ([Blei et al., 2017](#))—hence the name coordinate ascent variational inference (CAVI). Unlike the EM algorithm though, the CAVI algorithm does *not* guarantee an increase in the marginal log-likelihood at each step, nor does it guarantee convergence to the global maxima of the log-likelihood.

Further, the ELBO expression to be maximised is often not convex, which means the CAVI algorithm may terminate at local modes, for which there may be many. Note that the variational distribution with the higher ELBO value is the distribution that is closer, in terms of the KL divergence, to the true posterior distribution. In our experience, multiple random starts alleviates this issue for the I-probit model.

5.5 Post-estimation

5.6 Computational consideration

5.7 Examples

5.8 Conclusion

This work presents an extension of the normal I-prior methodology to fit categorical response models using probit link functions—a methodology we call the I-probit. The main motivation behind this work is to overcome the drawbacks of modelling probabilities using the normal I-prior model. We assumed latent variables that represent ‘class propensities’ exist, modelled these using a normal I-prior, and simply transformed them via a probit link function. In this way, all of the advantages of the I-prior methodology seen for the normal model are preserved for binary and multinomial regression as well.

The core of this work explores ways in which to overcome the intractable integral presented by the I-probit model in [\(5.7\)](#). Techniques such as quadrature methods, Laplace approximation and MCMC tend to fail, or are unsatisfactorily slow to accomplish. The main reason for this is the dimension of this integral, which is nm , and thus for large sample sizes and/or number of classes, is unfeasible with such methods. We turned to

sec:iprobit
eg

alg:caviipr
obit

Algorithm 1 CAVI for the I-probit model

```

1: procedure INITIALISATION
2:   Initialise  $\tilde{\mathbf{y}}^{*(0)}, \tilde{\mathbf{w}}^{(0)}, \tilde{\boldsymbol{\alpha}}^{(0)}, \tilde{\mathbf{H}}_{\eta^{(0)}}, \Psi$ 
3:    $t \leftarrow 0$ 
4: end procedure

5: while not converged do
6:   for  $i = 1, \dots, n$  do ▷ Update  $\mathbf{y}^*$ 
7:      $q^{(t+1)}(\mathbf{y}_{i\cdot}^*) \leftarrow {}^t\text{N}_m(\tilde{\boldsymbol{\alpha}}^{(t)} + \tilde{\mathbf{w}}^{(t)\top} \tilde{\mathbf{h}}_{\eta^{(t)}}(x_i), \Psi, \mathcal{C}_{y_i})$ 
8:      $\tilde{\mathbf{y}}_{i\cdot}^{*(t+1)} \leftarrow \text{E}_{q^{(t+1)}}[\mathbf{y}_{i\cdot}^*]$ 
9:   end for

10:   $\mathbf{V}_w^{(t+1)} \leftarrow ((\Psi \otimes \tilde{\mathbf{H}}_{\eta^{(t)}}^2) + (\Psi^{-1} \otimes \mathbf{I}_n))^{-1}$  ▷ Update  $\mathbf{w}$ 
11:   $\tilde{\mathbf{w}}^{(t+1)} \leftarrow \tilde{\mathbf{V}}_w^{(t+1)}(\Psi \otimes \tilde{\mathbf{H}}_{\eta^{(t)}}) \text{vec}(\tilde{\mathbf{y}}^{*(t+1)} - \mathbf{1}_n \tilde{\boldsymbol{\alpha}}^{(t)\top})$ 
12:   $q^{(t+1)}(\mathbf{w}) \leftarrow \text{N}_{nm}(\tilde{\mathbf{w}}^{(t+1)}, \mathbf{V}_w^{(t+1)})$ 

13:  Update  $q^{(t+1)}(\eta)$  as per Section 5.4.3 ▷ Update  $\eta$ 
14:  Sample  $\eta^{[1]}, \dots, \eta^{[T]} \sim q^{(t+1)}(\eta)$ 
15:   $\tilde{\mathbf{H}}_{\eta^{(t+1)}} \leftarrow \frac{1}{T} \sum_{i=1}^T \mathbf{H}_{\eta^{[i]}}$  and  $\tilde{\mathbf{H}}_{\eta^{(t+1)}}^2 \leftarrow \frac{1}{T} \sum_{i=1}^T \mathbf{H}_{\eta^{[i]}}^2$ 

16:   $\tilde{\boldsymbol{\alpha}}^{(t+1)} \leftarrow \frac{1}{n} \sum_{i=1}^n (\tilde{\mathbf{y}}_{i\cdot}^{*(t+1)} - \tilde{\mathbf{w}}^{(t+1)\top} \tilde{\mathbf{h}}_{\eta^{(t+1)}}(x_i))$  ▷ Update  $\boldsymbol{\alpha}$ 
17:   $q^{(t+1)}(\boldsymbol{\alpha}) \leftarrow \text{N}_m(\tilde{\boldsymbol{\alpha}}^{(t+1)}, \frac{1}{n} \Psi^{-1})$ 

18:  Calculate ELBO  $\mathcal{L}^{(t+1)}$ 
19:   $t \leftarrow t + 1$ 
20: end while

```

variational inference in the face of an intractable posterior density that hampers an EM algorithm, and the result is a sequential updating scheme, similar in time and storage requirements to the EM algorithm.

In terms of similarity to other works, the generalised additive models (GAMs) of [Hastie and Tibshirani \(1986\)](#) comes close. The setup of GAMs is near identical to the I-probit model, although estimation is done differently. GAMs do not assume smooth functions from any RKHS, but instead estimates the f 's using a local scoring method or a local likelihood method. Kernel methods for classification are extremely popular in computer science and machine learning; examples include support vector machines ([Schölkopf and Smola, 2002](#)) and Gaussian process classification ([Rasmussen and Williams, 2006](#)), with the latter being more closely related to the I-probit method. How-

ever, Gaussian process classification typically uses the logistic sigmoid function, and estimation most commonly performed using Laplace approximation, but other methods such as expectation propagation (Minka, 2001) and MCMC (Neal, 1999) have been explored as well. Variational inference for Gaussian process probit models have been studied by Girolami and Rogers (2006), with their work providing a close reference to the variational algorithm employed by us.

Suggestions for future work include:

1. **Estimation of Ψ .** A limitation we had to face in this work was to treat Ψ as fixed. This limitation was in part due to the non-conjugate nature of the variational density for Ψ . We believe the variational Bayes EM algorithm, which estimates maximum a posteriori values for the parameters, could alleviate this issue. This would bring the estimation procedure on par with the frequentist objective of maximum likelihood via the EM algorithm, albeit with the use of approximate posterior densities (see ???? for further discussions).
2. **Inclusion of class-specific covariates.** Throughout the chapter, we assumed that covariates were unit-specific, rather than class-specific. One such example is modelling the choice of travel mode between two destinations (car, coach, train or aeroplane) as a function of travel time. Clearly, travel time depends on the mode of transport. This would require a careful rethink of the appropriate RKHS/RKKS to which the regression function belongs: the regression on the latent propensities could be extended as such:

$$y_{ij}^* = \alpha_j + f_j(x_i) + e(z_{ij})$$

and $f_j \in \mathcal{F}_{\mathcal{X}}$, the RKHS with kernel $h : (\mathcal{X} \times \mathcal{M})^2 \rightarrow \mathbb{R}$ defined by $\delta_{jj'} h(x, x')$, and $e \in \mathcal{F}_{\mathcal{Z}}$, the RKHS of functions of the form $e : \{z_{ij} | i = 1, \dots, n, j = 1, \dots, m\} \rightarrow \mathbb{R}$. An I-prior would then be applied as usual, but the implications on the estimation would need to be considered as well.

3. **Improving computational efficiency.** The $O(n^3m)$ time requirement for estimating I-probit models hinder its use towards large-data applications. In a limited study, we did not obtain reliable improvements using low-rank approximations of the kernel matrix such as the Nyström method. The key to improving computational efficiency could lie in sparse variational methods, a suggestion that was made to improve normal I-prior models as well.




Figure 5.5: Time taken to complete a single variational inference iteration for varying sample sizes and number of classes m . The solid line represents actual timings, while the dotted lines are linear extrapolations.

5.9 Miscellanea

Appendix

5.10 Derivation of the CAVI algorithm

Let $\mathcal{Z} = \{\mathbf{y}^*, \mathbf{w}, \alpha, \eta, \Psi\}$. Approximate the posterior for \mathcal{Z} by a mean-field variational distribution

$$\begin{aligned} p(\mathbf{y}^*, \mathbf{w}, \alpha, \eta, \Psi | \mathbf{y}) &\approx q(\mathbf{y}^*)q(\mathbf{w})q(\alpha)q(\eta)q(\Psi) \\ &= \prod_{i=1}^n q(\mathbf{y}_i^*)q(\mathbf{w})q(\alpha)q(\eta)q(\Psi). \end{aligned}$$

The first line is by assumption, while the second line follows from an induced factorisation on the latent propensities, as we will see later. If needed, we also assume that $q(\eta)$ factorises into its constituents components. Recall that, for each $\xi \in \mathcal{Z}$, the optimal mean-field variational density \tilde{q} for ξ satisfies

$$\log \tilde{q}(\xi) = \mathbb{E}_{-\xi}[\log p(\mathbf{y}, \mathcal{Z})] + \text{const.} \quad (??)$$

Write $\mathbf{f} = \mathbf{H}_\eta \mathbf{w} \in \mathbb{R}^{n \times m}$. The joint likelihood $p(\mathbf{y}, \mathcal{Z})$ is given by

$$\begin{aligned} p(\mathbf{y}, \mathcal{Z}) &= p(\mathbf{y}|\mathcal{Z})p(\mathcal{Z}) \\ &= p(\mathbf{y}|\mathbf{y}^*)p(\mathbf{y}^*|\boldsymbol{\alpha}, \mathbf{w}, \eta, \boldsymbol{\Psi})p(\mathbf{w}|\boldsymbol{\Psi})p(\eta)p(\boldsymbol{\Psi})p(\boldsymbol{\alpha}). \end{aligned}$$

For reference, the relevant distributions are listed below.

- $p(\mathbf{y}|\mathbf{y}^*)$. For each observation $i \in \{1, \dots, n\}$, given the corresponding latent propensities $\mathbf{y}_i^* = (y_{i1}^*, \dots, y_{im}^*)$, the distribution for y_i is a degenerate distribution which depends on the j 'th component of \mathbf{y}_i^* being largest, where the value observed for y_i was j . Since each of the y_i 's are independent, everything is multiplicative.

$$p(\mathbf{y}|\mathbf{y}^*) = \prod_{i=1}^n \prod_{j=1}^m p_{ij}^{[y_i=j]} = \prod_{i=1}^n \prod_{j=1}^m \mathbb{1}[y_{ij}^* = \max_k y_{ik}^*] \mathbb{1}[y_i=j].$$

- $p(\mathbf{y}^*|\boldsymbol{\alpha}, \mathbf{w}, \eta, \boldsymbol{\Psi})$. Given values for the parameters and I-prior random effects, the distribution of the latent propensities is matrix normal

$$\mathbf{y}^*|\boldsymbol{\alpha}, \mathbf{w}, \eta, \boldsymbol{\Psi} \sim \text{MN}_{n,m}(\mathbf{1}_n \boldsymbol{\alpha}^\top + \mathbf{H}_\eta \mathbf{w}, \mathbf{I}_n, \boldsymbol{\Psi}^{-1}).$$

Write $\boldsymbol{\mu} = \mathbf{1}_n \boldsymbol{\alpha}^\top + \mathbf{H}_\eta \mathbf{w}$. Its pdf is

$$\begin{aligned} p(\mathbf{y}^*|\boldsymbol{\alpha}, \mathbf{w}, \eta, \boldsymbol{\Psi}) &= \exp \left[-\frac{nm}{2} \log 2\pi + \frac{n}{2} \log |\boldsymbol{\Psi}| - \frac{1}{2} \text{tr}((\mathbf{y}^* - \boldsymbol{\mu}) \boldsymbol{\Psi} (\mathbf{y}^* - \boldsymbol{\mu})^\top) \right] \\ &= \exp \left[-\frac{nm}{2} \log 2\pi + \frac{n}{2} \log |\boldsymbol{\Psi}| - \frac{1}{2} \sum_{i=1}^n (\mathbf{y}_{i\cdot}^* - \boldsymbol{\mu}_{i\cdot})^\top \boldsymbol{\Psi} (\mathbf{y}_{i\cdot}^* - \boldsymbol{\mu}_{i\cdot}) \right], \end{aligned}$$

where $\mathbf{y}_i^* \in \mathbb{R}^m$ and $\boldsymbol{\mu}_i \in \mathbb{R}^m$ are the rows of \mathbf{y}^* and $\boldsymbol{\mu}$ respectively. The second line follows directly from the definition of the trace, but also emanates from the fact that \mathbf{y}_i^* are independent multivariate normal with mean $\boldsymbol{\mu}_i$ and variance $\boldsymbol{\Psi}^{-1}$.

- $p(\mathbf{w}|\boldsymbol{\Psi})$. The \mathbf{w} 's are normal random matrices $\mathbf{w} \sim \text{MN}_{n,m}(\mathbf{0}, \mathbf{I}_n, \boldsymbol{\Psi})$ with pdf

$$\begin{aligned} p(\mathbf{w}|\boldsymbol{\Psi}) &= \exp \left[-\frac{nm}{2} \log 2\pi - \frac{n}{2} \log |\boldsymbol{\Psi}| - \frac{1}{2} \text{tr}(\mathbf{w} \boldsymbol{\Psi}^{-1} \mathbf{w}^\top) \right] \\ &= \exp \left[-\frac{nm}{2} \log 2\pi - \frac{n}{2} \log |\boldsymbol{\Psi}| - \frac{1}{2} \sum_{i=1}^n \mathbf{w}_{i\cdot}^\top \boldsymbol{\Psi}^{-1} \mathbf{w}_{i\cdot} \right]. \end{aligned}$$

- **$p(\boldsymbol{\eta})$** . The most common scenario would be $\boldsymbol{\eta} = \{\lambda_1, \dots, \lambda_p\}$ only. In this case, choose independent normal priors for each $\lambda_k \sim \mathcal{N}(m_k, v_k)$, $k = 1, \dots, p$, whose pdf is

$$p(\boldsymbol{\eta}) = \prod_{k=1}^p \exp \left[-\frac{1}{2} \log 2\pi - \frac{1}{2} \log v_k - \frac{1}{2v_k} (\lambda_k - m_k)^2 \right].$$

An improper prior $p(\boldsymbol{\eta}) \propto \text{const.}$ can be used as well, and this is the same as letting $m_k \rightarrow 0$ and $v_k \rightarrow 0$. The resulting posterior will be proper. If $\boldsymbol{\eta}$ contains other parameters as well, such as the Hurst coefficient $\gamma \in (0, 1)$, SE lengthscale $l > 0$ or polynomial offset $c > 0$, then appropriate priors should be used to match the support of the parameter. Choices include $p(\gamma) = \mathbb{1}(\gamma \in (0, 1))$ and $l, c \sim \Gamma(a, b)$.

- **$p(\boldsymbol{\Psi})$** . Our analysis shows that regardless of prior choice of $\boldsymbol{\Psi}$, be it in the full or independent I-probit model, the posterior for $\boldsymbol{\Psi}$ will not be of a recognisable form. Without giving too much thought, assume an improper prior on $\boldsymbol{\Psi}$, i.e. $p(\boldsymbol{\Psi}) \propto \text{const.}$
- **$p(\boldsymbol{\alpha})$** . Choose independent normal priors for the intercept, $\alpha_j \sim \mathcal{N}(a_j, A_j)$ for $j = 1, \dots, m$. The pdf is

$$p(\boldsymbol{\alpha}) = \prod_{j=1}^m \exp \left[-\frac{1}{2} \log 2\pi - \frac{1}{2} \log A_j - \frac{1}{2A_j} (\alpha_j - a_j)^2 \right].$$

Remark 5.2. The priors on the parameters $\{\boldsymbol{\alpha}, \boldsymbol{\eta}\}$ can be set to very vague or even improper priors, and the resulting posterior will still yield a proper distribution. Using improper priors eases the algebra slightly. For the precision matrix $\boldsymbol{\Psi}$, it is best to stick with the Wishart prior to avoid positive-definite issues, unless the independent I-probit model is used, in which case Jeffreys' prior for the precisions $p(\sigma_j^{-2}) \propto \sigma_j^2$ is a convenient choice.

5.10.1 Derivation of $\tilde{q}(\mathbf{y}^*)$

The rows of \mathbf{y}^* are independent, and thus we can consider the variational density for each \mathbf{y}_i^* separately. Consider the case where y_i takes one particular value $j \in \{1, \dots, m\}$.

The mean-field density $q(\mathbf{y}_i^*)$ for each $i = 1, \dots, n$ is found to be

$$\begin{aligned} \log \tilde{q}(\mathbf{y}_i^*) &= \mathbb{1}[y_{ij}^* = \max_k y_{ik}^*] \mathbb{E}_{\mathcal{Z} \setminus \{\mathbf{y}^*\} \sim q} \left[-\frac{1}{2} (\mathbf{y}_i^* - \boldsymbol{\mu}_i)^\top \boldsymbol{\Psi} (\mathbf{y}_i^* - \boldsymbol{\mu}_i) \right] + \text{const.} \\ &= \mathbb{1}[y_{ij}^* = \max_k y_{ik}^*] \left[-\frac{1}{2} (\mathbf{y}_i^* - \tilde{\boldsymbol{\mu}}_i)^\top \tilde{\boldsymbol{\Psi}} (\mathbf{y}_i^* - \tilde{\boldsymbol{\mu}}_i) \right] + \text{const.} \quad (\star) \\ &\equiv \begin{cases} \phi(\mathbf{y}_i^* | \tilde{\boldsymbol{\mu}}_i, \tilde{\boldsymbol{\Psi}}) & \text{if } y_{ij}^* > y_{ik}^*, \forall k \neq j \\ 0 & \text{otherwise} \end{cases} \end{aligned}$$

where $\tilde{\boldsymbol{\mu}}_i = \mathbb{E} \boldsymbol{\alpha} + (\mathbb{E} \mathbf{H}_\eta \mathbb{E} \mathbf{w})_i$, and expectations are taken under the optimal mean-field distribution \tilde{q} . The distribution $q(\mathbf{y}_i^*)$ is a truncated m -variate normal distribution such that the j 'th component is always largest. Unfortunately, the expectation of this distribution cannot be found in closed-form, and must be approximated by techniques such as Monte Carlo integration. If, however, the independent I-probit model is used and $\tilde{\boldsymbol{\Psi}}$ is diagonal, then ?? provides a simplification.

Remark 5.3. In (\star) above, we needn't consider the second order terms in the expectations because they do not involve \mathbf{y}^* and can be absorbed into the constant. To see this,

$$\begin{aligned} \mathbb{E}[(\mathbf{y}_i^* - \boldsymbol{\mu}_i)^\top \boldsymbol{\Psi} (\mathbf{y}_i^* - \boldsymbol{\mu}_i)] &= \mathbb{E}[\mathbf{y}_i^{*\top} \boldsymbol{\Psi} \mathbf{y}_i^* + \boldsymbol{\mu}_i^\top \boldsymbol{\Psi} \boldsymbol{\mu}_i - 2\boldsymbol{\mu}_i^\top \boldsymbol{\Psi} \mathbf{y}_i^*] \\ &= \mathbf{y}_i^{*\top} \boldsymbol{\Psi} \mathbf{y}_i^* - 2 \mathbb{E}[\boldsymbol{\mu}_i^\top] \mathbb{E}[\boldsymbol{\Psi}] \mathbf{y}_i^* + \text{const.} \\ &= \mathbf{y}_i^{*\top} \boldsymbol{\Psi} \mathbf{y}_i^* - 2\tilde{\boldsymbol{\mu}}_i^\top \tilde{\boldsymbol{\Psi}} \mathbf{y}_i^* + \text{const.} \\ &= (\mathbf{y}_i^* - \tilde{\boldsymbol{\mu}}_i)^\top \tilde{\boldsymbol{\Psi}} (\mathbf{y}_i^* - \tilde{\boldsymbol{\mu}}_i) + \text{const.} \end{aligned}$$

We will see this occurring a lot later on and we shall take note of this fact.

5.10.2 Derivation of $\tilde{q}(\mathbf{w})$

The terms involving \mathbf{w} in ?? are the $p(\mathbf{y}^* | \boldsymbol{\alpha}, \mathbf{w}, \eta, \boldsymbol{\Psi})$ and $p(\mathbf{w} | \boldsymbol{\Psi})$ terms, and the rest are absorbed into the constant. The easiest way to derive $\tilde{q}(\mathbf{w})$ is to vectorise \mathbf{y}^* and \mathbf{w} . We know that

$$\begin{aligned} \text{vec } \mathbf{y}^* | \boldsymbol{\alpha}, \mathbf{w}, \eta, \boldsymbol{\Psi} &\sim N_{nm} \left(\text{vec}(\mathbf{1}_n \boldsymbol{\alpha}^\top + \mathbf{H}_\eta \mathbf{w}), \boldsymbol{\Psi}^{-1} \otimes \mathbf{I}_n \right) \\ &\text{and} \\ \text{vec } \mathbf{w} | \boldsymbol{\Psi} &\sim N_{nm}(\mathbf{0}, \boldsymbol{\Psi} \otimes \mathbf{I}_n) \end{aligned}$$

using properties of matrix normal distributions. We also use the fact that $\text{vec}(\mathbf{H}_\eta \mathbf{w}) = (\mathbf{I}_m \otimes \mathbf{H}_\eta) \text{vec } \mathbf{w}$. For simplicity, write $\bar{\mathbf{y}}^* = \text{vec}(\mathbf{y}^* - \mathbf{1}_n \boldsymbol{\alpha}^\top)$, and $\mathbf{M} = (\mathbf{I}_m \otimes \mathbf{H}_\eta)$. Thus,

$$\begin{aligned} \log \tilde{q}(\mathbf{w}) &= E_{\mathcal{Z} \setminus \{\mathbf{w}\} \sim q} \left[-\frac{1}{2} (\bar{\mathbf{y}}^* - \mathbf{M} \text{vec } \mathbf{w})^\top (\boldsymbol{\Psi}^{-1} \otimes \mathbf{I}_n)^{-1} (\bar{\mathbf{y}}^* - \mathbf{M} \text{vec } \mathbf{w}) \right] \\ &\quad + E_{\mathcal{Z} \setminus \{\mathbf{w}\} \sim q} \left[-\frac{1}{2} (\text{vec } \mathbf{w})^\top (\boldsymbol{\Psi} \otimes \mathbf{I}_n)^{-1} \text{vec } (\mathbf{w}) \right] + \text{const.} \\ &= -\frac{1}{2} E_{\mathcal{Z} \setminus \{\mathbf{w}\} \sim q} \left[(\text{vec } \mathbf{w})^\top \overbrace{(\mathbf{M}^\top (\boldsymbol{\Psi} \otimes \mathbf{I}_n) \mathbf{M} + (\boldsymbol{\Psi}^{-1} \otimes \mathbf{I}_n))}^{\mathbf{A}} \text{vec } (\mathbf{w}) \right] \\ &\quad + E_{\mathcal{Z} \setminus \{\mathbf{w}\} \sim q} \left[\overbrace{\bar{\mathbf{y}}^{*\top} (\boldsymbol{\Psi} \otimes \mathbf{I}_n) \mathbf{M}}^{\mathbf{a}^\top} \text{vec } (\mathbf{w}) \right] + \text{const.} \\ &= -\frac{1}{2} E_{\mathcal{Z} \setminus \{\mathbf{w}\} \sim q} \left[(\text{vec } \mathbf{w} - \mathbf{A}^{-1} \mathbf{a})^\top \mathbf{A} (\text{vec } \mathbf{w} - \mathbf{A}^{-1} \mathbf{a}) \right] + \text{const.} \end{aligned}$$

This is recognised as a multivariate normal of dimension nm with mean and precision given by $\text{vec } \tilde{\mathbf{w}} = E[\mathbf{A}^{-1} \mathbf{a}]$ and $\tilde{\mathbf{V}}_w^{-1} = E[\mathbf{A}]$ respectively. With a little algebra, we find that

$$\begin{aligned} \mathbf{V}_w^{-1} &= E_{\mathcal{Z} \setminus \{\mathbf{w}\} \sim q} [\mathbf{A}] \\ &= E_{\mathcal{Z} \setminus \{\mathbf{w}\} \sim q} \left[(\mathbf{I}_m \otimes \mathbf{H}_\eta)^\top (\boldsymbol{\Psi} \otimes \mathbf{I}_n) (\mathbf{I}_m \otimes \mathbf{H}_\eta) + (\boldsymbol{\Psi}^{-1} \otimes \mathbf{I}_n) \right] \\ &= E_{\mathcal{Z} \setminus \{\mathbf{w}\} \sim q} \left[(\boldsymbol{\Psi} \otimes \mathbf{H}_\eta^2) + (\boldsymbol{\Psi}^{-1} \otimes \mathbf{I}_n) \right] \\ &= (\tilde{\boldsymbol{\Psi}} \otimes \tilde{\mathbf{H}}_\eta^2) + (\tilde{\boldsymbol{\Psi}}^{-1} \otimes \mathbf{I}_n) \end{aligned}$$

and making a first-order approximation $(E \mathbf{A})^{-1} \approx E[\mathbf{A}^{-1}]^3$,

$$\begin{aligned} \text{vec } \tilde{\mathbf{w}} &= E_{\mathcal{Z} \setminus \{\mathbf{w}\} \sim q} [\mathbf{A}^{-1} \mathbf{a}] \\ &= \tilde{\mathbf{V}}_w E_{\mathcal{Z} \setminus \{\mathbf{w}\} \sim q} \left[(\mathbf{I}_m \otimes \mathbf{H}_\eta) (\boldsymbol{\Psi} \otimes \mathbf{I}_n) \text{vec } (\mathbf{y}^* - \mathbf{1}_n \boldsymbol{\alpha}^\top) \right] \\ &= \tilde{\mathbf{V}}_w E_{\mathcal{Z} \setminus \{\mathbf{w}\} \sim q} \left[(\boldsymbol{\Psi} \otimes \mathbf{H}_\eta) \text{vec } (\mathbf{y}^* - \mathbf{1}_n \boldsymbol{\alpha}^\top) \right] \\ &= \tilde{\mathbf{V}}_w (\tilde{\boldsymbol{\Psi}} \otimes \tilde{\mathbf{H}}_\eta) \text{vec } (\tilde{\mathbf{y}}^* - \mathbf{1}_n \tilde{\boldsymbol{\alpha}}^\top). \end{aligned}$$

Ideally, we do not want to work with the $nm \times nm$ matrix \mathbf{V}_w , since its inverse is expensive to compute. Refer to ?? for details.

³Groves and Rothenberg (1969) show that $E[\mathbf{A}^{-1}] = (E \mathbf{A})^{-1} + \mathbf{B}$, where \mathbf{B} is a positive-definite matrix. This approximation has been used also by Girolami and Rogers (2006) in their work.

In the case of the I-probit model, where $\Psi = \text{diag}(\psi_1, \dots, \psi_m)$, then the covariance matrix takes a simpler form. Specifically, it has the block diagonal structure:

$$\begin{aligned}\tilde{\mathbf{V}}_w &= \text{E} [\text{diag}(\psi_1, \dots, \psi_m) \otimes \mathbf{H}_\eta^2 + \text{diag}(\psi_1, \dots, \psi_m) \otimes \mathbf{I}_n]^{-1} \\ &= \text{diag} \left(\text{E} (\psi_1 \mathbf{H}_\eta^2 + \psi_1^{-1} \mathbf{I}_n)^{-1}, \dots, \text{E} (\psi_m \mathbf{H}_\eta^2 + \psi_m^{-1} \mathbf{I}_n)^{-1} \right) \\ &\approx \text{diag} \left((\tilde{\psi}_1 \tilde{\mathbf{H}}_\eta^2 + \tilde{\psi}_1^{-1} \mathbf{I}_n)^{-1}, \dots, (\tilde{\psi}_m \tilde{\mathbf{H}}_\eta^2 + \tilde{\psi}_m^{-1} \mathbf{I}_n)^{-1} \right) \\ &=: \text{diag}(\tilde{\mathbf{V}}_{w_1}, \dots, \tilde{\mathbf{V}}_{w_m}).\end{aligned}$$

The mean $\text{vec } \tilde{\mathbf{w}}$ is

$$\begin{aligned}\text{vec } \tilde{\mathbf{w}} &= \tilde{\mathbf{V}}_w (\text{diag}(\tilde{\psi}_1, \dots, \tilde{\psi}_m) \otimes \tilde{\mathbf{H}}_\eta) \text{vec}(\tilde{\mathbf{y}}^* - \mathbf{1}_n \tilde{\boldsymbol{\alpha}}^\top) \\ &= \text{diag}(\tilde{\mathbf{V}}_{w_1}, \dots, \tilde{\mathbf{V}}_{w_m}) \text{diag}(\tilde{\psi}_1 \tilde{\mathbf{H}}_\eta, \dots, \tilde{\psi}_m \tilde{\mathbf{H}}_\eta) \text{vec}(\tilde{\mathbf{y}}^* - \mathbf{1}_n \tilde{\boldsymbol{\alpha}}^\top) \\ &= \text{diag}(\tilde{\psi}_1 \tilde{\mathbf{V}}_{w_1} \tilde{\mathbf{H}}_\eta, \dots, \tilde{\psi}_m \tilde{\mathbf{V}}_{w_m} \tilde{\mathbf{H}}_\eta) (\tilde{\mathbf{y}}^* - \mathbf{1}_n \tilde{\boldsymbol{\alpha}}^\top) \\ &= \begin{pmatrix} \tilde{\mathbf{w}}_{\cdot 1} & \dots & \tilde{\mathbf{w}}_{\cdot m} \end{pmatrix}^\top \\ &= \begin{pmatrix} \tilde{\psi}_1 \tilde{\mathbf{V}}_{w_1} \tilde{\mathbf{H}}_\eta (\tilde{\mathbf{y}}_{\cdot 1}^* - \tilde{\alpha}_1 \mathbf{1}_n) & \dots & \tilde{\psi}_m \tilde{\mathbf{V}}_{w_m} \tilde{\mathbf{H}}_\eta (\tilde{\mathbf{y}}_{\cdot m}^* - \tilde{\alpha}_m \mathbf{1}_n) \end{pmatrix}^\top.\end{aligned}$$

Therefore, we can consider the distribution of $\mathbf{w} = (\mathbf{w}_{\cdot 1}, \dots, \mathbf{w}_{\cdot m})$ columnwise, and each are normally distributed with mean and variance

$$\tilde{\mathbf{w}}_{\cdot j} = \tilde{\sigma}_j^{-2} \tilde{\mathbf{V}}_{w_j} \tilde{\mathbf{H}}_\eta (\tilde{\mathbf{y}}_{\cdot j}^* - \tilde{\alpha}_j \mathbf{1}_n) \quad \text{and} \quad \tilde{\mathbf{V}}_{w_j} = (\tilde{\sigma}_j^{-2} \tilde{\mathbf{H}}_\eta^2 + \tilde{\sigma}_j^2 \mathbf{I}_n)^{-1}.$$

A quantity that we will be requiring time and again will be $\text{tr}(\mathbf{C} \text{E}[\mathbf{w}^\top \mathbf{D} \mathbf{w}])$, where $\mathbf{C} \in \mathbb{R}^{m \times m}$ and $\mathbf{D} \in \mathbb{R}^{n \times n}$ are both square and symmetric matrices. Using the definition of the trace directly, we get

$$\begin{aligned}\text{tr}(\mathbf{C} \text{E}[\mathbf{w}^\top \mathbf{D} \mathbf{w}]) &= \sum_{i,j=1}^m \mathbf{C}_{ij} \text{E}[\mathbf{w}^\top \mathbf{D} \mathbf{w}]_{ij} \\ &= \sum_{i,j=1}^m \mathbf{C}_{ij} \text{E}[\mathbf{w}_{\cdot i}^\top \mathbf{D} \mathbf{w}_{\cdot j}].\end{aligned} \tag{5.11}$$

{eq:trCEwDw}

The expectation of the univariate quantity $\mathbf{w}_{\cdot i}^\top \mathbf{D} \mathbf{w}_{\cdot j}$ is inspected below:

$$\begin{aligned} \mathbb{E}[\mathbf{w}_{\cdot i}^\top \mathbf{D} \mathbf{w}_{\cdot j}] &= \text{tr}(\mathbf{D} \mathbb{E}[\mathbf{w}_{\cdot j} \mathbf{w}_{\cdot i}^\top]) \\ &= \text{tr}(\mathbf{D}(\text{Cov}(\mathbf{w}_{\cdot j}, \mathbf{w}_{\cdot i}) + \mathbb{E}[\mathbf{w}_{\cdot j}] \mathbb{E}[\mathbf{w}_{\cdot i}]^\top)) \\ &= \text{tr}(\mathbf{D}(\mathbf{V}_w[i, j] + \tilde{\mathbf{w}}_{\cdot j} \tilde{\mathbf{w}}_{\cdot i}^\top)). \end{aligned}$$

where $\mathbf{V}_w[i, j] \in \mathbb{R}^{n \times n}$ refers to the (i, j) 'th submatrix block of \mathbf{V}_w . Of course, in the independent the I-probit model, this is equal to

$$\mathbf{V}_w[i, j] = \delta_{ij}(\psi_j \mathbf{H}_\eta^2 + \psi_j^{-1} \mathbf{I}_n)^{-1}$$

where δ is the Kronecker delta. Continuing on (5.11) leads us to

$$\text{tr}(\mathbf{C} \mathbb{E}[\mathbf{w}^\top \mathbf{D} \mathbf{w}]) = \sum_{i,j=1}^m \mathbf{C}_{ij} \left(\text{tr}(\mathbf{D}(\delta_{ij} \mathbf{V}_{w_j} + \tilde{\mathbf{w}}_{\cdot j} \tilde{\mathbf{w}}_{\cdot i}^\top)) \right).$$

If $\mathbf{C} = \text{diag}(c_1, \dots, c_m)$, then

$$\begin{aligned} \text{tr}(\mathbf{C} \mathbb{E}[\mathbf{w}^\top \mathbf{D} \mathbf{w}]) &= \sum_{j=1}^m c_j \left(\text{tr}(\mathbf{D} \tilde{\mathbf{V}}_{w_j}) + \tilde{\mathbf{w}}_{\cdot j}^\top \mathbf{D} \tilde{\mathbf{w}}_{\cdot j} \right) \\ &= \sum_{j=1}^m c_j \text{tr}(\mathbf{D}(\tilde{\mathbf{V}}_{w_j} + \tilde{\mathbf{w}}_{\cdot j} \tilde{\mathbf{w}}_{\cdot j}^\top)) \end{aligned}$$

5.10.3 Derivation of $\tilde{q}(\eta)$

By looking at only the terms involving η in ??, we deduce that \tilde{q} for η satisfies

$$\begin{aligned} \log \tilde{q}(\eta) &= -\frac{1}{2} \text{tr} \mathbb{E}_{\mathcal{Z} \setminus \{\eta\} \sim q} \left[(\mathbf{y}^* - \mathbf{1}_n \boldsymbol{\alpha}^\top - \mathbf{H}_\eta \mathbf{w}) \boldsymbol{\Psi} (\mathbf{y}^* - \mathbf{1}_n \boldsymbol{\alpha}^\top - \mathbf{H}_\eta \mathbf{w})^\top \right] + \log p(\eta) \\ &\quad + \text{const.} \\ &= -\frac{1}{2} \text{tr} \mathbb{E}_{\mathcal{Z} \setminus \{\eta\} \sim q} \left(\boldsymbol{\Psi} \mathbf{w}^\top \mathbf{H}_\eta^2 \mathbf{w} - 2 \boldsymbol{\Psi} \mathbf{w}^\top \mathbf{H}_\eta (\mathbf{y}^* - \boldsymbol{\alpha}) \right) + \log p(\eta) + \text{const.} \\ &= -\frac{1}{2} \text{tr} \left(\tilde{\boldsymbol{\Psi}} \mathbb{E}[\mathbf{w}^\top \mathbf{H}_\eta^2 \mathbf{w}] - 2 \tilde{\boldsymbol{\Psi}} \tilde{\mathbf{w}}^\top \mathbf{H}_\eta (\tilde{\mathbf{y}}^* - \tilde{\boldsymbol{\alpha}}) \right) + \log p(\eta) + \text{const.} \end{aligned}$$

with some appropriate prior $p(\eta)$. In general, this does not have a recognisable form in η , especially when it is not linearly dependent on the kernel matrix. This happens when considering parameters other than the scales of the RKHSs. Our interest would

be to obtain $\tilde{\mathbf{H}}_\eta := \mathbb{E}_{\eta \sim q} \mathbf{H}_\eta$ and $\tilde{\mathbf{H}}_\eta^2 := \mathbb{E}_{\eta \sim q} \mathbf{H}_\eta^2$. We use a Metropolis random-walk algorithm to obtain these quantities, as detailed in the algorithm below.

Algorithm 2 Metropolis random-walk to sample η

- 1: **inputs** $\tilde{\alpha}$, $\tilde{\mathbf{w}}$, $\tilde{\Psi}$, and s Metropolis sampling s.d.
- 2: **initialise** $\eta^{(0)} \in \mathbb{R}^q$ and $t \leftarrow 0$
- 3: **for** $t = 1, \dots, T$ **do**
- 4: Draw $\eta^* \sim N_q(\eta^{(t)}, s^2)$
- 5: Accept/reject proposal state, i.e.

$$\eta^{(t+1)} \leftarrow \begin{cases} \eta^* & \text{if } u \sim \text{Unif}(0, 1) < \pi_{\text{acc}} \\ \eta^{(t)} & \text{otherwise} \end{cases}$$

where

$$\pi_{\text{acc}} = \min \left(1, \exp \left(\log \tilde{q}(\eta^*) - \log \tilde{q}(\eta^{(t)}) \right) \right).$$

- 6: **end for**
 - 7: $\tilde{\mathbf{H}}_\eta \leftarrow \frac{1}{T} \sum_{i=1}^T \mathbf{H}_{\eta^{(i)}}$ and $\tilde{\mathbf{H}}_\eta^2 \leftarrow \frac{1}{T} \sum_{i=1}^T \mathbf{H}_{\eta^{(i)}}^2$
-

Now consider the case where $\eta = \{\lambda_1, \dots, \lambda_p\}$ (RKHS scale parameters only), and the scenario described in the exponential family EM algorithm of [Section 4.3.3](#) applies. In particular, for $k = 1, \dots, p$, we can decompose the kernel matrix as $\mathbf{H}_\eta = \lambda_k \mathbf{R}_k + \mathbf{S}_k$ and its square as $\mathbf{H}_\eta^2 = \lambda_k^2 \mathbf{R}_k^2 + \lambda_k \mathbf{U}_k + \mathbf{S}_k^2$. Then, for $j = 1, \dots, m$, assuming each of

the $q(\lambda_k)$ densities are independent of each other, we find that

$$\begin{aligned}
 \log \tilde{q}(\lambda_k) &= \mathbb{E}_{\mathcal{Z} \setminus \{\eta\} \sim q} \left[-\frac{1}{2} \text{tr} \left((\mathbf{y}^* - \boldsymbol{\mu}) \boldsymbol{\Psi} (\mathbf{y}^* - \boldsymbol{\mu})^\top \right) \right] - \frac{1}{2v_k^2} (\lambda_k - m_k)^2 + \text{const.} \\
 &= -\frac{1}{2} \text{tr} \mathbb{E}_{\mathcal{Z} \setminus \{\eta\} \sim q} \left[\boldsymbol{\Psi} \mathbf{w}^\top \mathbf{H}_\eta^2 \mathbf{w} - 2\boldsymbol{\Psi} (\mathbf{y}^* - \mathbf{1} \boldsymbol{\alpha}^\top)^\top \mathbf{H}_\eta \mathbf{w} \right] \\
 &\quad - \frac{1}{2v_k^2} (\lambda_k - m_k)^2 + \text{const.} \\
 &= -\frac{1}{2} \text{tr} \mathbb{E}_{\mathcal{Z} \setminus \{\eta\} \sim q} \left[\boldsymbol{\Psi} \mathbf{w}^\top (\lambda_k^2 \mathbf{R}_k^2 + \lambda_k \mathbf{U}_k) \mathbf{w} - 2\boldsymbol{\Psi} (\mathbf{y}^* - \mathbf{1} \boldsymbol{\alpha}^\top)^\top (\lambda_k \mathbf{R}_k) \mathbf{w} \right] \\
 &\quad - \frac{1}{2v_k^2} (\lambda_k^2 - 2m_k \lambda_k) + \text{const.} \\
 &= -\frac{1}{2} \text{tr} \mathbb{E}_{\mathcal{Z} \setminus \{\eta\} \sim q} \left[\lambda_k^2 \boldsymbol{\Psi} \mathbf{w}^\top \mathbf{R}_k^2 \mathbf{w} - 2\lambda_k \left(\boldsymbol{\Psi} (\mathbf{y}^* - \mathbf{1} \boldsymbol{\alpha}^\top)^\top \mathbf{R}_k \mathbf{w} - \frac{1}{2} \boldsymbol{\Psi} \mathbf{w}^\top \mathbf{U}_k \mathbf{w} \right) \right] \\
 &\quad - \frac{1}{2} \left(\frac{1}{v_k^2} \lambda_k^2 - 2 \frac{m_k}{v_k^2} \lambda_k \right) + \text{const.} \\
 &= -\frac{1}{2} \left[\lambda_k^2 \overbrace{\left(\text{tr}(\tilde{\boldsymbol{\Psi}} \mathbb{E}[\mathbf{w}^\top \mathbf{R}_k^2 \mathbf{w}]) \right)}^{c_k} + v_k^{-2} \right] \\
 &\quad - 2\lambda_k \overbrace{\left(\text{tr} \left(\tilde{\boldsymbol{\Psi}} (\tilde{\mathbf{y}}^* - \mathbf{1}_n \tilde{\boldsymbol{\alpha}}^\top)^\top \mathbf{R}_k \tilde{\mathbf{w}} - \frac{1}{2} \tilde{\boldsymbol{\Psi}} \mathbb{E}[\mathbf{w}^\top \mathbf{U}_k \mathbf{w}] \right) + m_k v_k^{-2} \right)}^{d_k} \Big]
 \end{aligned}$$

By completing the squares, we recognise this is as the kernel of a univariate normal density. Specifically, $\lambda_k \sim \mathcal{N}(d_k/c_k, 1/c_k)$. The quantity $\tilde{\mathbf{H}}_\eta$ can be obtained by substituting $\lambda_k \mapsto \mathbb{E}_{\lambda_k \sim q}[\lambda_k]$ in the expression $\mathbf{H}_\eta = \lambda_k \mathbf{R}_k + \mathbf{S}_k$. However, in the calculation of $\tilde{\mathbf{H}}_\eta^2$ using $\lambda_k^2 \mathbf{R}_k^2 + \lambda_k \mathbf{U}_k + \mathbf{S}_k^2$, we must replace occurrences of λ_k^2 with $\mathbb{E}_{\lambda_k \sim q}[\lambda_k]^2 + \text{Var}_{\lambda_k \sim q}[\lambda_k]$. This can be cumbersome, so if felt necessary, use the approximation $\lambda_k^2 \mapsto \mathbb{E}_{\lambda_k \sim q}[\lambda_k]^2$ instead.

Example 5.1. Suppose $k = 1$, and we only have λ to estimate. Then, $\mathbf{H}_\eta = \lambda \mathbf{H}$, $\mathbf{R}_k = \mathbf{H}$, $\mathbf{R}_k^2 = \mathbf{H}^2$, and $\mathbf{U}_k = \mathbf{0}$. Suppose also we use an improper prior $\lambda_k \propto \text{const.}$, which is the same as having $v_k^2 \rightarrow 0$ and $m_k v_k^{-2} \rightarrow 0$. The mean field distribution for λ is then

$$\lambda \sim \mathcal{N} \left(\frac{\text{tr}(\tilde{\boldsymbol{\Psi}}(\tilde{\mathbf{y}}^* - \mathbf{1} \tilde{\boldsymbol{\alpha}}^\top)^\top \mathbf{H} \tilde{\mathbf{w}})}{\text{tr}(\tilde{\boldsymbol{\Psi}} \mathbb{E}[\mathbf{w}^\top \mathbf{H}^2 \mathbf{w}])}, \frac{1}{\text{tr}(\tilde{\boldsymbol{\Psi}} \mathbb{E}[\mathbf{w}^\top \mathbf{H}^2 \mathbf{w}])} \right)$$

Further, if $\tilde{\boldsymbol{\Psi}} = \tilde{\psi} \mathbf{I}_m$, then

$$\lambda \sim \mathcal{N} \left(\frac{\sum_{j=1}^m (\tilde{\mathbf{y}}_{\cdot j}^* - \tilde{\alpha}_j \mathbf{1})^\top \mathbf{H} \tilde{\mathbf{w}}_{\cdot j}}{\sum_{j=1}^m \text{tr}(\mathbf{H}^2 \mathbb{E}[\mathbf{w}_{\cdot j} \mathbf{w}_{\cdot j}^\top])}, \frac{1}{\sum_{j=1}^m \text{tr}(\mathbf{H}^2 \mathbb{E}[\mathbf{w}_{\cdot j} \mathbf{w}_{\cdot j}^\top])} \right)$$

which bears a resemblance to the exponential family EM algorithm solutions described in Chapter 4. Now, $\tilde{\mathbf{H}}_\eta = \mathbb{E}[\lambda \mathbf{H}] = \tilde{\lambda} \mathbf{H}$, and $\tilde{\mathbf{H}}_\eta^2 = \mathbb{E}[\lambda^2 \mathbf{H}^2] = (\text{Var } \lambda + \tilde{\lambda}^2) \mathbf{H}^2$.

5.10.4 Derivation of $\tilde{q}(\Psi)$

We find that $q(\Psi)$ satisfies

$$\begin{aligned} \log q(\Psi) &= \mathbb{E}_{\mathcal{Z} \setminus \{\Psi\} \sim q} \left[-\frac{1}{2} \text{tr}((\mathbf{y}^* - \boldsymbol{\mu})^\top (\mathbf{y}^* - \boldsymbol{\mu}) \Psi) - \frac{1}{2} \text{tr}(\mathbf{w}^\top \mathbf{w} \Psi^{-1}) \right] \\ &\quad + \log p(\Psi) + \text{const.} \\ &= -\frac{1}{2} \text{tr} \left(\overbrace{(\mathbb{E}[(\mathbf{y}^* - \boldsymbol{\mu})^\top (\mathbf{y}^* - \boldsymbol{\mu})])}^{\mathbf{G}_1} \Psi + \overbrace{\mathbb{E}[\mathbf{w}^\top \mathbf{w}]}^{\mathbf{G}_2} \Psi^{-1} \right) \\ &\quad + \log p(\Psi) + \text{const.} \end{aligned}$$

This seems to be the pdf of $\text{Wis}(\mathbf{G} + \mathbf{G}_1, g)$ plus the pdf of a distribution which almost resembles an inverse Wishart pdf. Unfortunately, the properties such as its moments and entropy are unknown.

The matrix \mathbf{G}_1 is

$$\begin{aligned} \mathbf{G}_1 &= \mathbb{E}[(\mathbf{y}^* - \boldsymbol{\mu})^\top (\mathbf{y}^* - \boldsymbol{\mu})] \\ &= \mathbb{E}[\mathbf{y}^{*\top} \mathbf{y}^* + \boldsymbol{\alpha} \mathbf{1}_n^\top \mathbf{1}_n \boldsymbol{\alpha}^\top + \mathbf{w}^\top \mathbf{H}_\eta^2 \mathbf{w} - 2\mathbf{y}^{*\top} \mathbf{1}_n \boldsymbol{\alpha}^\top - 2\mathbf{y}^{*\top} \mathbf{H}_\eta \mathbf{w} - 2\boldsymbol{\alpha} \mathbf{1}_n^\top \mathbf{H}_\eta \mathbf{w}] \\ &= \mathbb{E}[\mathbf{y}^{*\top} \mathbf{y}^*] + n \mathbb{E}[\boldsymbol{\alpha} \boldsymbol{\alpha}^\top] + \mathbb{E}[\mathbf{w}^\top \mathbf{H}_\eta \mathbf{w}] - 2(\tilde{\mathbf{y}}^{*\top} \mathbf{1}_n \tilde{\boldsymbol{\alpha}}^\top + \tilde{\mathbf{y}}^{*\top} \tilde{\mathbf{H}}_\eta \tilde{\mathbf{w}} + \tilde{\boldsymbol{\alpha}} \mathbf{1}_n^\top \tilde{\mathbf{H}}_\eta \tilde{\mathbf{w}}), \end{aligned}$$

and this involves second order moments of a conically truncated multivariate normal distribution, which needs to be obtained via simulation. Meanwhile,

$$\begin{aligned} \mathbf{G}_{2,ij} &= \mathbb{E}[\mathbf{w}^\top \mathbf{w}]_{ij} \\ &= \mathbb{E}[\mathbf{w}_{\cdot i}^\top \mathbf{w}_{\cdot j}] \\ &= \tilde{\mathbf{V}}_w[i, j] + \tilde{\mathbf{w}}_{\cdot i}^\top \tilde{\mathbf{w}}_{\cdot j}. \end{aligned}$$

In the case of the independent I-probit model, we use a gamma prior on each of the precisions in the diagonal entries of $\Psi = \text{diag}(\psi_1, \dots, \psi_m)$. Then, the variational density

for each ψ_j is found to be

$$\begin{aligned} \log q(\psi_j) &= \mathbb{E}_{\mathcal{Z} \setminus \{\Psi\} \sim q} \left[\frac{n}{2} \log(\psi_1 \cdots \psi_m) - \frac{1}{2} \sum_{j=1}^m \sum_{i=1}^n \psi_j (\mathbf{y}_{ij}^* - \boldsymbol{\mu}_{ij})^2 \right] \\ &\quad + \mathbb{E}_{\mathcal{Z} \setminus \{\Psi\} \sim q} \left[-\frac{n}{2} \log(\psi_1 \cdots \psi_m) - \frac{1}{2} \sum_{j=1}^m \sum_{i=1}^n \psi_j^{-1} \mathbf{w}_{ij}^2 \right] \\ &\quad + \sum_{j=1}^m ((s_j - 1) \log \psi_j - r_j \psi_j) + \text{const.} \\ &= (s_j - 1) \log \psi_j - \psi_j \left(\frac{1}{2} \mathbb{E} \|\mathbf{y}_{\cdot,j}^* - \boldsymbol{\mu}_{\cdot,j}\|^2 + r_j \right) \\ &\quad - \psi_j^{-1} \left(\frac{1}{2} \mathbb{E} \|\mathbf{w}_{\cdot,j}\|^2 \right) + \text{const.} \end{aligned}$$

which again, is a pdf of an unknown distribution. However, its posterior mode can be computed. Write $a = -\left(\frac{1}{2} \mathbb{E} \|\mathbf{y}_{\cdot,j}^* - \boldsymbol{\mu}_{\cdot,j}\|^2 + r_j\right)$, $b = s_j - 1$, and $c = \left(\frac{1}{2} \mathbb{E} \|\mathbf{w}_{\cdot,j}\|^2\right)$. Then,

$$\frac{\partial}{\partial \psi_j} \log q(\psi_j) = \frac{\partial}{\partial \psi_j} (a\psi_j + b \log \psi_j - c\psi_j^{-1}) = a + b\psi_j^{-1} + c\psi_j^{-2}$$

equated to zero means solving a quadratic equation in ψ_j . Suppose that $p(\psi_j) \propto \text{const.}$, then $s_j = 1$ and $r_j = 0$ so $\tilde{\psi}_j$ can be solved directly to be

$$\hat{\psi}_j = \sqrt{\frac{\mathbb{E} \|\mathbf{y}_{\cdot,j}^* - \boldsymbol{\mu}_{\cdot,j}\|^2}{\mathbb{E} \|\mathbf{w}_{\cdot,j}\|^2}}.$$

If the posterior mean is close to its mode, then $\hat{\psi}_j$ is a good approximation for $\tilde{\psi}_j$.

To calculate $\mathbb{E} \|\mathbf{y}_{\cdot,j}^* - \boldsymbol{\mu}_{\cdot,j}\|^2 = \mathbb{E} \sum_{i=1}^n (\mathbf{y}_{ij}^* - \mu_{ij})^2$, one first needs $\mathbb{E} (y_{ij}^* - \alpha_j - \mathbf{w}_{\cdot,j}^\top \mathbf{h}_\eta(x_i))^2$. This, in itself, presents a challenge to compute analytically, because it requires, among other things, the second moments $\mathbb{E} y_{ij}^{*2}$ and $\mathbb{E} [\mathbf{w}_{\cdot,j}^\top \mathbf{h}_\eta(x_i) \mathbf{h}_\eta(x_i)^\top \mathbf{w}_{\cdot,j}]$. Although not entirely accurate, it is simpler to use the approximation

$$\mathbb{E} \|\mathbf{y}_{\cdot,j}^* - \boldsymbol{\mu}_{\cdot,j}\|^2 \approx \|\tilde{\mathbf{y}}_{\cdot,j}^* - \tilde{\boldsymbol{\mu}}_{\cdot,j}\|^2.$$

(see note ?? on page ??). Also, we have $\mathbf{w}_{\cdot,j} \sim N_n(\tilde{\mathbf{w}}_{\cdot,j}, \tilde{\mathbf{V}}_{w_j})$, and so $\mathbb{E} \|\mathbf{w}_{\cdot,j}\|^2 = \text{tr}(\tilde{\mathbf{V}}_{w_j} + \tilde{\mathbf{w}}_{\cdot,j} \tilde{\mathbf{w}}_{\cdot,j}^\top)$.

5.10.5 Derivation of $\tilde{q}(\boldsymbol{\alpha})$

Let $\mathbf{A} = \text{diag}(A_1, \dots, A_m)$ and $\mathbf{a} = (a_1, \dots, a_m)^\top$. The terms involving α_j in ?? are

$$\begin{aligned} \log q(\boldsymbol{\alpha}) &= \mathbb{E}_{\mathcal{Z} \setminus \{\boldsymbol{\alpha}\} \sim q} \left[-\frac{1}{2} \sum_{i=1}^n (\mathbf{y}_{i.}^* - \boldsymbol{\alpha} - \mathbf{w}^\top \mathbf{h}_\eta(x_i))^\top \boldsymbol{\Psi} (\mathbf{y}_{i.}^* - \boldsymbol{\alpha} - \mathbf{w}^\top \mathbf{h}_\eta(x_i)) \right] \\ &\quad - \frac{1}{2} (\boldsymbol{\alpha} - \mathbf{a})^\top \mathbf{A}^{-1} (\boldsymbol{\alpha} - \mathbf{a}) + \text{const.} \\ &= -\frac{1}{2} \left[\boldsymbol{\alpha}^\top \overbrace{(n\boldsymbol{\Psi} + \mathbf{A}^{-1})}^{\tilde{\mathbf{A}}} \boldsymbol{\alpha} - 2 \overbrace{\left(\sum_{i=1}^n \boldsymbol{\Psi} (\tilde{\mathbf{y}}_{i.}^* - \tilde{\mathbf{w}}^\top \tilde{\mathbf{h}}_\eta(x_i)) + \mathbf{A}^{-1} \mathbf{a} \right)}^{\tilde{\mathbf{a}}} \boldsymbol{\alpha} \right] \end{aligned}$$

which implies a normal mean-field distribution for $\boldsymbol{\alpha}$ whose mean and variance are $\tilde{\boldsymbol{\alpha}} = \tilde{\mathbf{A}}^{-1} \tilde{\mathbf{a}}$ and $\tilde{\mathbf{A}}^{-1}$ respectively. If $\boldsymbol{\Psi}$ is diagonal, the components of $\boldsymbol{\alpha}$ would be independent.

As a remark, due to identifiability, only $m - 1$ of these intercept are estimable. We can either put a constraint that one of the intercepts is fixed at zero, or the sum of the intercepts equals zero. The latter constraint is implemented in this thesis, and this is realised by estimating all the intercepts and then centring them.

Bibliography

- | | |
|-------------------------|---|
| albert1993bayesian | Albert, James H and Siddhartha Chib (1993). “Bayesian analysis of binary and polychotomous response data”. In: <i>Journal of the American statistical Association</i> 88.422, pp. 669–679. |
| blei2017variational | Blei, David M, Alp Kucukelbir, and Jon D McAuliffe (2017). “Variational inference: A review for statisticians”. In: <i>Journal of the American Statistical Association</i> just-accepted. |
| bunch1991estimability | Bunch, David S (1991). “Estimability in the multinomial probit model”. In: <i>Transportation Research Part B: Methodological</i> 25.1, pp. 1–12. |
| dansie1985parameter | Dansie, BR (1985). “Parameter estimability in the multinomial probit model”. In: <i>Transportation Research Part B: Methodological</i> 19.6, pp. 526–528. |
| girolami2006variational | Girolami, Mark and Simon Rogers (2006). “Variational Bayesian Multinomial Probit Regression with Gaussian Process Priors”. In: <i>Neural Computation</i> 18.8, pp. 1790–1817. |
| groves1969note | Groves, Theodore and Thomas Rothenberg (1969). “A note on the expected value of an inverse matrix”. In: <i>Biometrika</i> 56.3, pp. 690–691. |
| hastie1986 | Hastie, Trevor and Robert Tibshirani (Aug. 1986). “Generalized Additive Models”. In: <i>Statist. Sci.</i> 1.3, pp. 297–310. DOI: 10.1214/ss/1177013604 . URL: https://doi.org/10.1214/ss/1177013604 . |
| kass1995bayes | Kass, Robert E and Adrian E Raftery (1995). “Bayes factors”. In: <i>Journal of the american statistical association</i> 90.430, pp. 773–795. |
| Keane1992 | Keane, Michael P. (1992). “A Note on Identification in the Multinomial Probit Model”. In: <i>Journal of Business & Economic Statistics</i> 10.2, pp. 193–200. ISSN: 0735-0015. |

DOI: [10.2307/1391677](https://doi.org/10.2307/1391677). URL: <http://www.jstor.org/stable/1391677>
<http://www.jstor.org/stable/pdfplus/1391677.pdf?acceptTC=true>.

kuss2005ass
essing

Kuss, Malte and Carl Edward Rasmussen (2005). “Assessing approximate inference for binary Gaussian process classification”. In: *Journal of machine learning research* 6.Oct, pp. 1679–1704.

mccullagh19
89

McCullagh, P. and John A. Nelder (1989). *Generalized Linear Models*. 2nd. Chapman & Hall/CRC Press.

minka2001ex
pectation

Minka, Thomas P (2001). “Expectation propagation for approximate Bayesian inference”. In: *Proceedings of the Seventeenth conference on Uncertainty in artificial intelligence*. Morgan Kaufmann Publishers Inc., pp. 362–369.

neal1999

Neal, Radford M. (1999). “Regression and Classification using Gaussian Process Priors”. In: *Bayesian Statistics*. Ed. by J M Bernardo, J O Berger, A P Dawid, and A F M Smith. Vol. 6. Oxford University Press. (with discussion), pp. 475–501.

rasmussen20
06gaussian

Rasmussen, Carl Edward and Christopher K I Williams (2006). *Gaussian Processes for Machine Learning*. The MIT Press.

scholkopf20
02learning

Schölkopf, Bernhard and Alexander J Smola (2002). *Learning with Kernels: Support Vector Machines, Regularization, Optimization, and Beyond*. MIT Press.

train2009di
crete

Train, Kenneth E (2009). *Discrete choice methods with simulation*. Cambridge university press.